

References: Chapter 17 of Geron's book. For 1-Dim plots, Keras tutorial :

<https://www.tensorflow.org/tutorials/generative/autoencoder>

This file trains a modified VAE (with a different sampling layer and a different loss function) with the instances of the normal digits in the training data.

Then, it measures the reconstruction loss for the digits in the test data.

The reconstruction loss for the instances of the abnormal digits in the test data is higher.

A threshold is determined based on the distribution of the reconstruction losses of the normal training data (threshold = mean + 2.5*std of this distribution).

Then, if the reconstruction loss of a digit in the test data is higher than this threshold, it is classified as abnormal.

By comparing with the known labels of test data (with T for normal digit(s) and F for abnormal digit(s)), the confusion matrix and the accuracy is calculated.

```
import sklearn
import tensorflow as tf
from tensorflow import keras
import numpy as np

import matplotlib as mpl
import matplotlib.pyplot as plt
mpl.rc('axes', labelsize=14)
mpl.rc('xtick', labelsize=12)
mpl.rc('ytick', labelsize=12)

from sklearn.metrics import accuracy_score, precision_score, recall_score, confusion_matrix
from sklearn.model_selection import train_test_split
```

Loading the MNIST data and forming arrays of the normal training data, the validation data (normal and abnormal) and the test data (normal and abnormal)

```
#Labels
# 0 T-shirt/top
# 1 Trouser
# 2 Pullover
# 3 Dress
# 4 Coat
# 5 Sandal
```

```

# 6 Shirt
# 7 Sneaker
# 8 Bag
# 9 Ankle boot

nl1 = 6
nl2 = 6
abn1 = 7
abn2 = 7

(x_train_0, y_train_0), (x_test, y_test) = keras.datasets.fashion_mnist.load_data()

x_train_0 = x_train_0.astype(np.float32) / 255
x_test = x_test.astype(np.float32) / 255

train_size = x_train_0.shape[0] * 9 // 10

x_train, x_valid, y_train, y_valid = train_test_split(x_train_0, y_train_0, train_size = train_size)

normal_data = x_train[(y_train == nl1) | (y_train == nl2)]           # Normal training data (Normal)
normal_labels = y_train[(y_train == nl1) | (y_train == nl2)]

valid_data = x_valid[(y_valid == abn1) | (y_valid == abn2) | (y_valid == nl1) | (y_valid == nl2)]
valid_labels = y_valid[(y_valid == abn1) | (y_valid == abn2) | (y_valid == nl1) | (y_valid == nl2)]

test_data = x_test[(y_test == abn1) | (y_test == abn2) | (y_test == nl1) | (y_test == nl2)]
test_labels = y_test[(y_test == abn1) | (y_test == abn2) | (y_test == nl1) | (y_test == nl2)]

test_labels_T_F = np.where((test_labels == nl1) | (test_labels == nl2), True, False)
# Array of T and F, T where test digits are normal and F where test digits are abnormal

valid_labels_T_F = np.where((valid_labels == nl1) | (valid_labels == nl2), True, False)
# Array of T and F, T where test digits are normal and F where test digits are abnormal

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-32768/29515 [=====] - 0s 0us/step
40960/29515 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-26427392/26421880 [=====] - 0s 0us/step
26435584/26421880 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-16384/5148 [=====]
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-4423680/4422102 [=====] - 0s 0us/step
4431872/4422102 [=====] - 0s 0us/step

normal_data.shape, normal_labels.shape, valid_data.shape, valid_labels.shape, test_data.shape
((5402, 28, 28), (5402,), (1218, 28, 28), (1218,), (2000, 28, 28), (2000,))

```

```
normal_test_data = test_data[(test_labels == nl1) | (test_labels == nl2)] # The no
abnormal_test_data = test_data[(test_labels == abn1) | (test_labels == abn2)] # The ab
normal_test_labels = test_labels[(test_labels == nl1) | (test_labels == nl2)] # Their
abnormal_test_labels = test_labels[(test_labels == abn1) | (test_labels == abn2)] # Their
```

```
normal_test_data.shape, abnormal_test_data.shape
```

```
((1000, 28, 28), (1000, 28, 28))
```

```
normal_valid_data = valid_data[(valid_labels == nl1) | (valid_labels == nl2)] # Th
abnormal_valid_data = valid_data[(valid_labels == abn1) | (valid_labels == abn2)] # Th
normal_valid_labels = valid_labels[(valid_labels == nl1) | (valid_labels == nl2)] # Th
abnormal_valid_labels = valid_labels[(valid_labels == abn1) | (valid_labels == abn2)] # Th
```

```
normal_valid_data.shape, abnormal_valid_data.shape
```

```
((598, 28, 28), (620, 28, 28))
```

▼ Building and training the network

```
K = keras.backend
# def rounded_accuracy(y_true, y_pred):
#     return keras.metrics.binary_accuracy(tf.round(y_true), tf.round(y_pred))

# Modified sampling layer with the addition of mean_2, log_var_2, and fraction p, with
# the appropriate change in the reparametrization trick to do stochastic
# sampling from the superposition of the two MVN distributions, while allowing
# the 5 parallel layers containing the means and stds of the two MVNs and the fractions p's
# for each dimension to be trained via backpropogation of the error signal.
class Sampling(keras.layers.Layer):
    def call(self, inputs):
        mean_1, log_var_1, mean_2, log_var_2, p = inputs
        return (K.random_normal(tf.shape(log_var_1)) * K.exp(log_var_1 / 2) + mean_1)*p + (K.

# For details please see Geron's book.
codings_size = 16 # The number of dimensions of the two MVN distributions in the sampling la

inputs = keras.layers.Input(shape=[28, 28])
z = keras.layers.Flatten()(inputs)
z = keras.layers.Dense(256, activation="selu")(z)
z = keras.layers.Dense(128, activation="selu")(z)
z = keras.layers.Dense(64, activation="selu")(z)

# Adding output nodes (parallel layers) at the end of the encoder for means
# and standard deviations of a second Multivariate Normal (MVN) distribution
# in the dimensions of the coding size (here 32). In each of the dimensions,
```

```

# this first MVN is multiplied by a fraction p and added to the second MVN
# multiplied by 1 - p in each dimension.
# final distribution = p * first MVN + (1 - p) * second MVN
# Another parallel layer (set of nodes) is added to keep and train the fractions p's
# in each dimension
codings_mean_1 = keras.layers.Dense(codings_size)(z)
codings_log_var_1 = keras.layers.Dense(codings_size)(z)
codings_mean_2 = keras.layers.Dense(codings_size)(z)
codings_log_var_2 = keras.layers.Dense(codings_size)(z)
codings_p = keras.layers.Dense(1, activation='sigmoid')(z)

# Modified sampling layer at the end of the encoder
codings = Sampling()([codings_mean_1, codings_log_var_1, codings_mean_2, codings_log_var_2, c
variational_encoder = keras.models.Model(
    inputs=[inputs], outputs=[codings_mean_1, codings_log_var_1, codings_mean_2, codings_log_
decoder_inputs = keras.layers.Input(shape=[codings_size])
x = keras.layers.Dense(64, activation="selu")(decoder_inputs)
x = keras.layers.Dense(128, activation="selu")(x)
x = keras.layers.Dense(256, activation="selu")(x)
x = keras.layers.Dense(28 * 28, activation="sigmoid")(x)
outputs = keras.layers.Reshape([28, 28])(x)
variational_decoder = keras.models.Model(inputs=[decoder_inputs], outputs=[outputs])

_, _, _, _, codings = variational_encoder(inputs)
reconstructions = variational_decoder(codings)
variational_ae = keras.models.Model(inputs=[inputs], outputs=[reconstructions])

# New latent loss function that will be added to the reconstruction binary cross-entropy loss
# The whole network (Encoder, sampling layer, and decoder) will train to minimize this loss
p_mean = K.mean(codings_p)
array1 = p_mean*(codings_log_var_1 - K.exp(codings_log_var_1) - K.square(codings_mean_1))
array2 = (1-p_mean)*(codings_log_var_2 - K.exp(codings_log_var_2) - K.square(codings_mean_2))
sum1 = K.sum(1 + array1, axis=-1)
sum2 = K.sum(1 + array2, axis=-1)

latent_loss = -0.5 * (sum1 + sum2)

latent_loss = latent_loss * 16

# Add the latent loss to the reconstruction loss
variational_ae.add_loss(K.mean(latent_loss) / 784.)

# For the reconstruction loss binary cross-entropy loss is used (same as regular VAE).
# For details please see Chapter 17 of Geron's book (Stacked AE and VAE sections)
variational_ae.compile(loss="binary_crossentropy", optimizer="rmsprop")

checkpoint_cb = keras.callbacks.ModelCheckpoint("modVAE_latent_times_16_model", monitor="val_
history = variational_ae.fit(normal_data, normal_data, epochs=100, batch_size=128, callbacks=

```

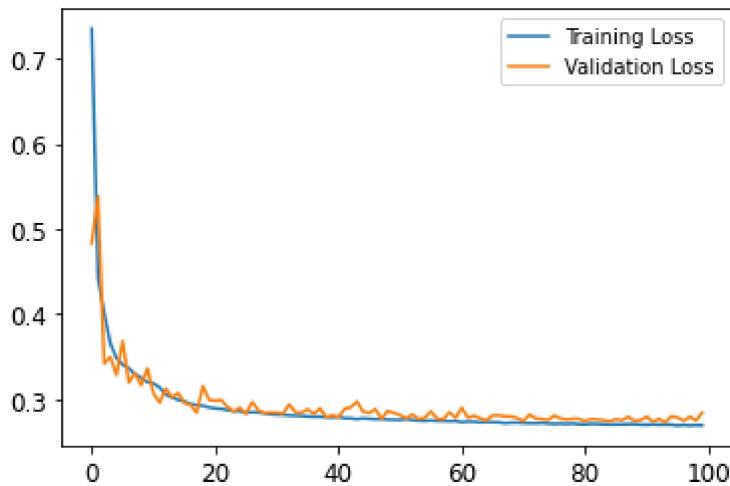
```
validation_data=(normal_valid_data, normal_valid_data), shuffle=
```

```
Epoch 74/100
43/43 [=====] - 0s 8ms/step - loss: 0.2726 - val_loss: 0.277
Epoch 75/100
43/43 [=====] - 0s 8ms/step - loss: 0.2719 - val_loss: 0.276
Epoch 76/100
43/43 [=====] - 0s 8ms/step - loss: 0.2716 - val_loss: 0.281
Epoch 77/100
43/43 [=====] - 0s 9ms/step - loss: 0.2718 - val_loss: 0.277
Epoch 78/100
43/43 [=====] - 0s 8ms/step - loss: 0.2714 - val_loss: 0.276
Epoch 79/100
43/43 [=====] - 0s 9ms/step - loss: 0.2718 - val_loss: 0.277
Epoch 80/100
43/43 [=====] - 0s 8ms/step - loss: 0.2712 - val_loss: 0.277
Epoch 81/100
43/43 [=====] - ETA: 0s - loss: 0.2710INFO:tensorflow:Assets
43/43 [=====] - 3s 74ms/step - loss: 0.2710 - val_loss: 0.274
Epoch 82/100
43/43 [=====] - 0s 8ms/step - loss: 0.2715 - val_loss: 0.277
Epoch 83/100
43/43 [=====] - 0s 9ms/step - loss: 0.2714 - val_loss: 0.276
Epoch 84/100
43/43 [=====] - 0s 8ms/step - loss: 0.2710 - val_loss: 0.275
Epoch 85/100
37/43 [=====>.....] - ETA: 0s - loss: 0.2705INFO:tensorflow:Assets
43/43 [=====] - 3s 66ms/step - loss: 0.2707 - val_loss: 0.274
Epoch 86/100
43/43 [=====] - 0s 8ms/step - loss: 0.2708 - val_loss: 0.276
Epoch 87/100
43/43 [=====] - 0s 8ms/step - loss: 0.2708 - val_loss: 0.275
Epoch 88/100
43/43 [=====] - 0s 9ms/step - loss: 0.2711 - val_loss: 0.279
Epoch 89/100
43/43 [=====] - 0s 9ms/step - loss: 0.2712 - val_loss: 0.276
Epoch 90/100
43/43 [=====] - 0s 9ms/step - loss: 0.2705 - val_loss: 0.276
Epoch 91/100
43/43 [=====] - 0s 8ms/step - loss: 0.2704 - val_loss: 0.279
Epoch 92/100
42/43 [=====>.] - ETA: 0s - loss: 0.2709INFO:tensorflow:Assets
43/43 [=====] - 3s 73ms/step - loss: 0.2708 - val_loss: 0.274
Epoch 93/100
43/43 [=====] - 0s 8ms/step - loss: 0.2705 - val_loss: 0.277
Epoch 94/100
42/43 [=====>.] - ETA: 0s - loss: 0.2707INFO:tensorflow:Assets
43/43 [=====] - 3s 67ms/step - loss: 0.2706 - val_loss: 0.274
Epoch 95/100
43/43 [=====] - 0s 8ms/step - loss: 0.2705 - val_loss: 0.280
Epoch 96/100
43/43 [=====] - 0s 9ms/step - loss: 0.2694 - val_loss: 0.278
Epoch 97/100
43/43 [=====] - 0s 9ms/step - loss: 0.2702 - val_loss: 0.275
Epoch 98/100
43/43 [=====] - 0s 9ms/step - loss: 0.2705 - val_loss: 0.278
```

```
43/43 [=====] - 0s 8ms/step - loss: 0.2700 - val_loss: 0.275
Epoch 100/100
43/43 [-----] - 0s 8ms/step - loss: 0.2607 - val_loss: 0.284
```

```
plt.plot(history.history["loss"], label="Training Loss")
plt.plot(history.history["val_loss"], label="Validation Loss")
plt.legend()
```

```
<matplotlib.legend.Legend at 0x7fe971656e50>
```



```
model = variational_ae
model.summary(expand_nested=True, show_trainable=True)
```

tf.math.square (TFOpLambda)	(None, 16)	0	'tf.math.exp[0][0]'
tf.math.subtract_3 (TFOpLambda)	(None, 16)	0	'dense_3[0][0]'
)			'tf.math.exp_1[0][0]
tf.math.square_1 (TFOpLambda)	(None, 16)	0	'dense_5[0][0]'
tf.math.subtract_1 (TFOpLambda)	(None, 16)	0	'tf.math.subtract[0
)			'tf.math.square[0][0]
tf.math.subtract_2 (TFOpLambda)	()	0	'tf.math.reduce_mean'
)			
tf.math.subtract_4 (TFOpLambda)	(None, 16)	0	'tf.math.subtract_3
)			'tf.math.square_1[0
tf.math.multiply (TFOpLambda)	(None, 16)	0	'tf.math.reduce_mean'
			'tf.math.subtract_1
tf.math.multiply_1 (TFOpLambda)	(None, 16)	0	'tf.math.subtract_2
)			'tf.math.subtract_4
tf._operators_.add (TFOpLamb	(None, 16)	0	'tf.math.multiply[0

da)

tf.__operators__.add_1 (TFOpLambda (None, 16) mbda)	0	['tf.math.multiply_1
tf.math.reduce_sum (TFOpLambda (None,))	0	['tf.__operators__.a
tf.math.reduce_sum_1 (TFOpLambda (None,) da)	0	['tf.__operators__.a]
tf.__operators__.add_2 (TFOpLambda (None,) mbda)	0	['tf.math.reduce_sum 'tf.math.reduce_sum_
tf.math.multiply_2 (TFOpLambda (None,))	0	['tf.__operators__.a]
tf.math.multiply_3 (TFOpLambda (None,))	0	['tf.math.multiply_2
tf.math.reduce_mean_1 (TFOpLambda () bda)	0	['tf.math.multiply_3
tf.math.truediv (TFOpLambda) ()	0	['tf.math.reduce_mean
add_loss (AddLoss) ()	0	['tf.math.truediv[0]

=====

Total params: 490,257

Trainable params: 490,257

Non-trainable params: 0

```
model_encoder = variational_encoder
# model_encoder.summary(expand_nested=True, show_trainable=True)

model_decoder = variational_decoder
# model_decoder.summary(expand_nested=True, show_trainable=True)

model_layers = np.array(model.layers)
n_layers = model_layers.shape[0]
# np.concatenate((np.arange(n_layers).reshape(n_layers,1), model_layers.reshape(n_layers,1)),
```

The original and reconstructed images for the first 30 instances of the normal training data, validation data, normal validation data, abnormal validation data, test data, normal test data, and abnormal test data

```
def plot_image(image):
```

```
plt.imshow(image, cmap="binary")
plt.axis("off")

def show_reconstructions(model, images=test_data, n_images=5):
    reconstructions = model.predict(images[:n_images])
    fig = plt.figure(figsize=(n_images * 1.5, 3))
    for image_index in range(n_images):
        plt.subplot(2, n_images, 1 + image_index)
        plot_image(images[image_index])
        plt.subplot(2, n_images, 1 + n_images + image_index)
        plot_image(reconstructions[image_index])

show_reconstructions(variational_ae, normal_data, 30)
plt.show()
```



```
show_reconstructions(variational_ae, valid_data, 30)
plt.show()
```



```
show_reconstructions(variational_ae, normal_valid_data, 30)
plt.show()
```



```
show_reconstructions(variational_ae, abnormal_valid_data, 30)
plt.show()
```



```
show_reconstructions(variational_ae, test_data, 30)
plt.show()
```



```
show_reconstructions(variational_ae, normal_test_data, 30)
plt.show()
```



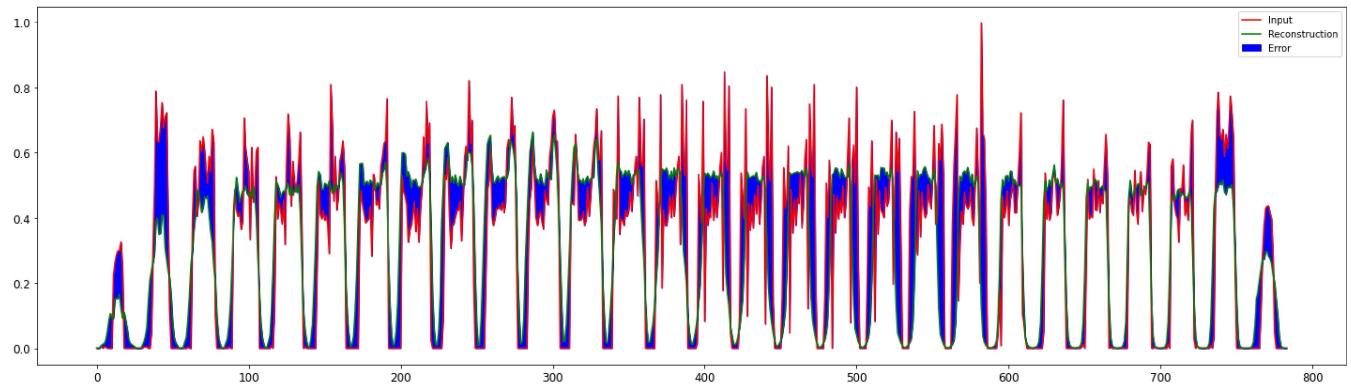
```
show_reconstructions(variational_ae, abnormal_test_data, 30)
plt.show()
```



1-Dim plot of pixels of the first normal test data

```
reconstructions_nl_test = variational_ae.predict(normal_test_data)
```

```
plt.figure(figsize=(25,7))
plt.plot(normal_test_data[0].ravel(), 'r')
plt.plot(reconstructions_nl_test[0].ravel(), 'g')
plt.fill_between(np.arange(28*28), reconstructions_nl_test[0].ravel(), normal_test_data[0].ravel())
plt.legend(labels=["Input", "Reconstruction", "Error"])
plt.show()
```

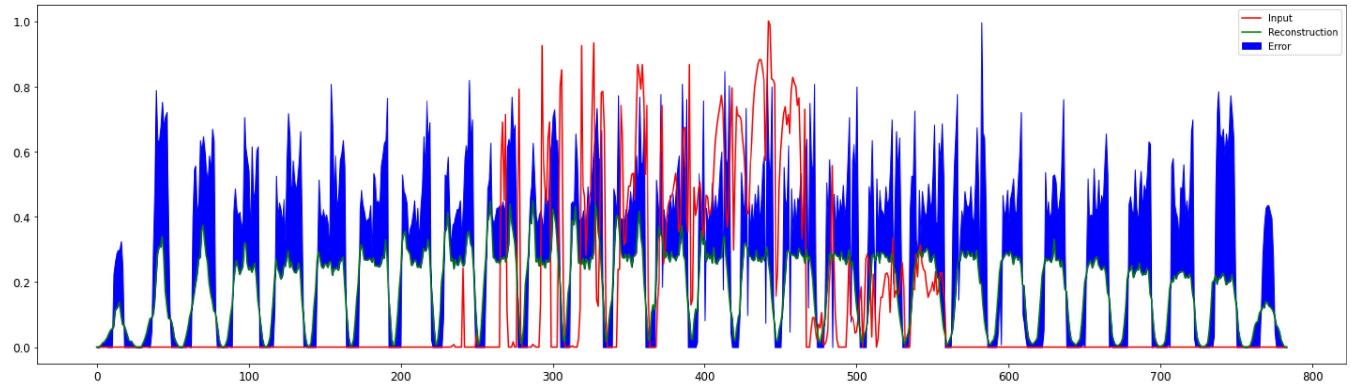


1-Dim plot of pixels of the first abnormal test data

```
reconstructions_abn_test = variational_ae.predict(abnormal_test_data)
```

```
plt.figure(figsize=(25,7))
plt.plot(abnormal_test_data[0].ravel(), 'r')
plt.plot(reconstructions_abn_test[0].ravel(), 'g')
plt.fill_between(np.arange(28*28), reconstructions_abn_test[0].ravel(), normal_test_data[0].ravel())
plt.show()
```

```
plt.legend(labels=["Input", "Reconstruction", "Error"])
plt.show()
```



▼ Distributions of the reconstruction losses and the calculation of the threshold.

Distribution of the reconstruction losses of the normal training data

```
reconstructions = variational_ae.predict(normal_data)
train_loss = tf.keras.losses.mae(reconstructions.reshape(-1, 784), normal_data.reshape(-1, 784))
plt.figure(figsize=(10,5))
plt.hist(train_loss[None,:], bins=100)
threshold1 = np.mean(train_loss) + 2.5*np.std(train_loss)
plt.axvline(threshold1,c='g')
plt.xlabel("MAE reconstruction loss of the normal training data")
plt.ylabel("No of examples")
plt.show()
```

```
print("Mean: ", np.mean(train_loss))
print("Std: ", np.std(train_loss))
```

```
Mean:  0.121271096
Std:  0.041283876
```

≤ 100



```
threshold_train_mean_2_5_std = np.mean(train_loss) + 2.5*np.std(train_loss)
print("Threshold based on the mean of the training data MAE reconstruction losses + 2.5 std: ")
```

```
Threshold based on the mean of the training data MAE reconstruction losses + 2.5 std:  0.181271096
```

≤ 20

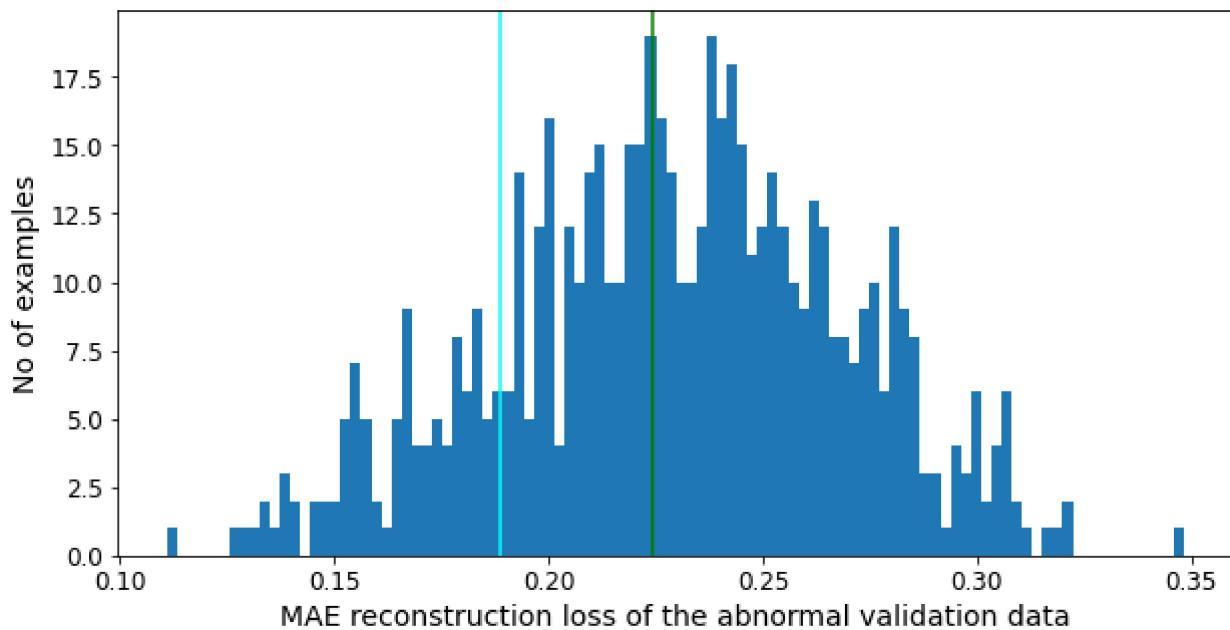
20

\geq

```
threshold1 = threshold_train_mean_2_5_std
```

Distribution of the reconstruction losses of the abnormal validation data

```
reconstructions = variational_ae.predict(abnormal_valid_data)
abn_valid_loss = tf.keras.losses.mae(reconstructions.reshape(-1,784), abnormal_valid_data.reshape(-1,784))
plt.figure(figsize=(10,5))
plt.hist(abn_valid_loss[None, :], bins=100)
threshold2 = np.mean(abn_valid_loss) - np.std(abn_valid_loss)
plt.axvline(threshold2,c='cyan')
plt.axvline(threshold1,c='g')
plt.xlabel("MAE reconstruction loss of the abnormal validation data")
plt.ylabel("No of examples")
plt.show()
```



```
abnormal_valid_mean_loss = np.mean(abn_valid_loss)
```

```
abnormal_valid_mean_loss , np.std(abn_valid_loss)
```

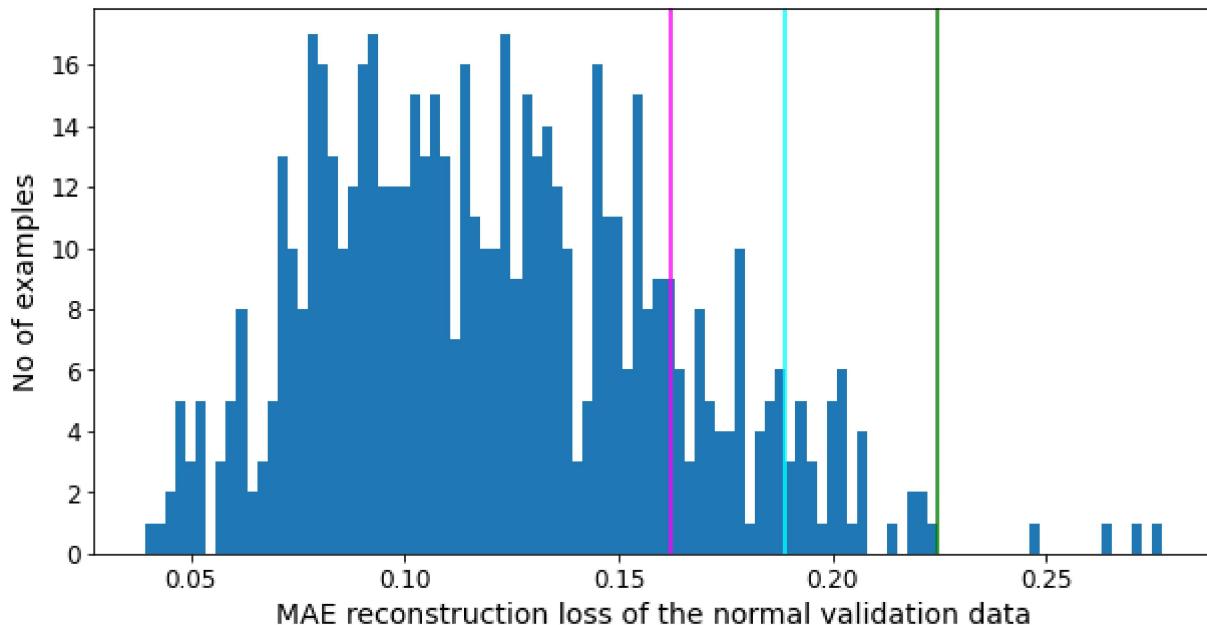
```
(0.229011, 0.040306848)
```

```
threshold2 = abnormal_valid_mean_loss - np.std(abn_valid_loss)
print("Threshold2: ", threshold2)
```

```
Threshold2: 0.18870415
```

Distribution of the reconstruction losses of the normal validation data

```
reconstructions = variational_ae.predict(normal_valid_data)
nl_valid_loss = tf.keras.losses.mae(reconstructions.reshape(-1,784), normal_valid_data.reshape(-1,784))
plt.figure(figsize=(10,5))
plt.hist(nl_valid_loss[None, :], bins=100)
threshold3 = np.mean(nl_valid_loss) + np.std(nl_valid_loss)
plt.axvline(threshold3, c='magenta')
plt.axvline(threshold2, c='cyan')
plt.axvline(threshold1, c='g')
plt.xlabel("MAE reconstruction loss of the normal validation data")
plt.ylabel("No of examples")
plt.show()
```



```
normal_valid_mean_loss = np.mean(nl_valid_loss)
```

```
normal_valid_mean_loss , np.std(nl_valid_loss)
```

```
(0.12159238, 0.04074322)
```

```
threshold3 = normal_valid_mean_loss + np.std(nl_valid_loss)
print("Threshold3: ", threshold3)
```

Threshold3: 0.1623356

Calculation of a preliminary threshold based on (threshold2 + threshold3) / 2 = Average of (mean + std of the distribution of the reconstruction losses of the normal validation data) and (mean - std of the distribution of the reconstruction losses of the abnormal validation data)

```
Avg_of_threshold_2_3 = (threshold2 + threshold3)/2
print("Average of threshold 2 and 3: ", Avg_of_threshold_2_3)
```

Average of threshold 2 and 3: 0.1755198836326599

```
threshold4 = Avg_of_threshold_2_3
```

Calculation of the threshold that gives the best accuracy on the validation data and set this as the threshold.

```
def predict(model, data, threshold):
    reconstructions = model.predict(data)
    loss = tf.keras.losses.mae(reconstructions.reshape(-1, 784), data.reshape(-1, 784))
    return tf.math.less(loss, threshold)
```

```
increment = (abnormal_valid_mean_loss - normal_valid_mean_loss)/100
thresholds = np.arange(normal_valid_mean_loss, abnormal_valid_mean_loss, increment)
thrs_size = thresholds.shape[0]
accuracies = np.zeros(thrs_size)
for i in range(thrs_size):
    preds = predict(variational_ae, valid_data, thresholds[i])
    accuracies[i] = accuracy_score(preds, valid_labels_T_F)
argmax = np.argmax(accuracies)
valid_data_best_threshold = thresholds[argmax]
print("The best threshold based on validation data: ", valid_data_best_threshold)
```

The best threshold based on validation data: 0.16348564155399803

```
thr_acc = np.zeros((thrs_size, 2))
thr_acc[:, 0] = thresholds
thr_acc[:, 1] = accuracies
thr_acc[argmax-2:argmax+3]
```

```
array([[0.16133727, 0.89737274],
       [0.16241146, 0.89737274],
```

```
[0.16348564, 0.8998358 ],
[0.16455983, 0.88669951],
[0.16563401, 0.89162562]])
```

```
threshold5 = valid_data_best_threshold
```

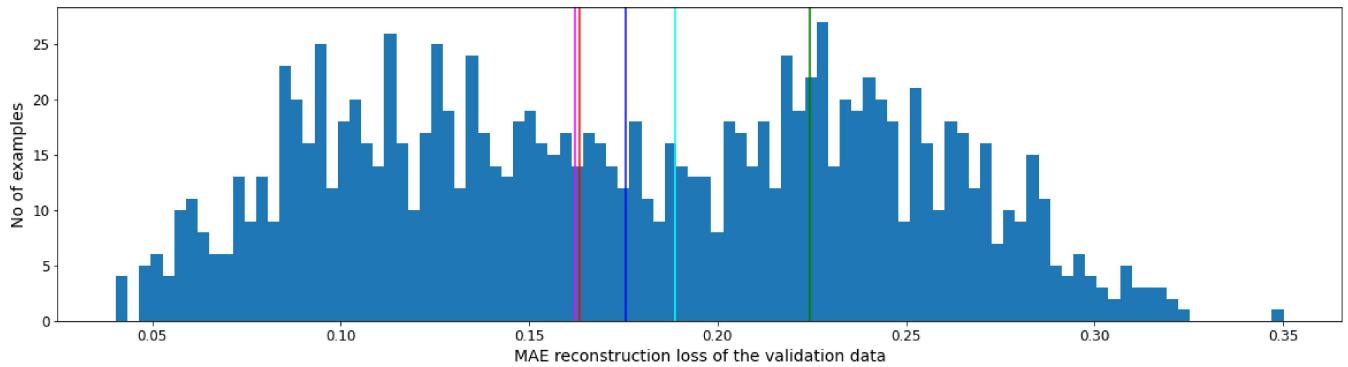
```
threshold = threshold5
```

Distribution of the reconstruction losses of all the validation data (normal and abnormal)

The blue line is threshold4 (= the average of threshold3 [magenta] and threshold2 [cyan]).

The red line is the threshold that gives the best accuracy for the validation data.

```
reconstructions = variational_ae.predict(valid_data)
valid_loss = tf.keras.losses.mae(reconstructions.reshape(-1,784), valid_data.reshape(-1,784))
plt.figure(figsize=(20,5))
plt.hist(valid_loss[None, :], bins=100)
plt.axvline(threshold, c='r')
plt.axvline(threshold4, c='b')
plt.axvline(threshold2, c='cyan')
plt.axvline(threshold3, c='magenta')
plt.axvline(threshold1, c='green')
plt.xlabel("MAE reconstruction loss of the validation data")
plt.ylabel("No of examples")
plt.show()
```

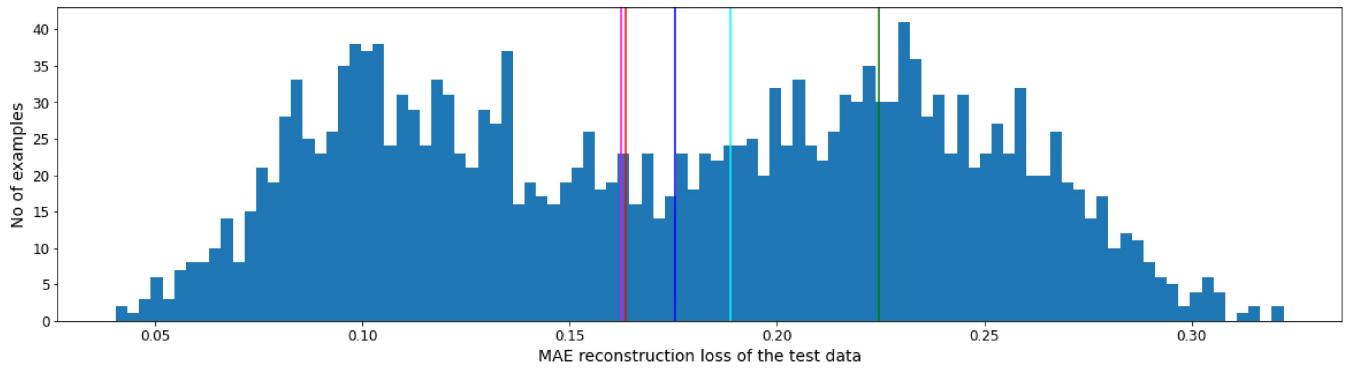


Distribution of the reconstruction losses of the test data (normal and abnormal)

The blue line is threshold4 (= the average of threshold3 [magenta] and threshold2 [cyan]).

The red line is the threshold that gives the best accuracy for the validation data.

```
reconstructions = variational_ae.predict(test_data)
test_loss = tf.keras.losses.mae(reconstructions.reshape(-1,784), test_data.reshape(-1,784))
plt.figure(figsize=(20,5))
plt.hist(test_loss[None, :], bins=100)
plt.axvline(threshold, c='r')
plt.axvline(threshold4, c='b')
plt.axvline(threshold2, c='cyan')
plt.axvline(threshold3, c='magenta')
plt.axvline(threshold1, c='green')
plt.xlabel("MAE reconstruction loss of the test data")
plt.ylabel("No of examples")
plt.show()
```



Mean and standard deviation of reconstruction losses for normal and abnormal test data

```
reconstructions = variational_ae.predict(normal_test_data)
nl_test_loss = tf.keras.losses.mae(reconstructions.reshape(-1,784), normal_test_data.reshape(-1,784))
np.mean(nl_test_loss) , np.std(nl_test_loss)

(0.12066414, 0.040881366)
```

```
reconstructions = variational_ae.predict(abnormal_test_data)
abn_test_loss = tf.keras.losses.mae(reconstructions.reshape(-1,784), abnormal_test_data.reshape(-1,784))
np.mean(abn_test_loss) , np.std(abn_test_loss)

(0.22774687, 0.036684614)
```

Calculation of the accuracy and the confusion matrix on the test data with threshold set based on the best threshold from the validation data

```

# def predict(model, data, threshold):
#     reconstructions = model.predict(data)
#     loss = tf.keras.losses.mae(reconstructions.reshape(-1, 784), data.reshape(-1, 784))
#     return tf.math.less(loss, threshold)

def print_stats(predictions, labels):
    cf = confusion_matrix(labels, predictions)
    print("Confusion Matrix: \n prediction: F      T ")
    print("          {}  {}".format(preds == False).shape[0], preds == True].sh
    print(" label: F  [[{}  {}]]  {}".format(cf[0,0], cf[0,1], test_labels_T_F[test_labels_T_
    print("      T  [[{}  {}]]  {}".format(cf[1,0], cf[1,1], test_labels_T_F[test_labels_T_
    print("Accuracy = {}".format(accuracy_score(labels, predictions)))
    print("Normal Test Data Mean = {}".format(np.mean(nl_test_loss)))
    print("Normal Test Data Standard Deviation = {}".format(np.std(nl_test_loss)))
    print("Abnormal Test Data Mean = {}".format(np.mean(abn_test_loss)))
    print("Abnormal Test Data Standard Deviation = {}".format(np.std(abn_test_loss)))
    print("Precision = {}".format(precision_score(labels, predictions)))
    print("Recall = {}".format(recall_score(labels, predictions)))
    print(accuracy_score(labels, predictions))
    print(np.mean(nl_test_loss))
    print(np.std(nl_test_loss))
    print(np.mean(abn_test_loss))
    print(np.std(abn_test_loss))
    print(precision_score(labels, predictions))
    print(recall_score(labels, predictions))
    print(accuracy_score(labels, predictions), np.mean(nl_test_loss), np.std(nl_test_loss), np.
        precision_score(labels, predictions), recall_score(labels, predictions))

preds = predict(variational_ae, test_data, threshold)
print_stats(preds, test_labels_T_F)

```

↳ Confusion Matrix:

		F	T
	F	1113	887
label: F	[[961	39]	1000
T	[152	848]]	1000

Accuracy = 0.9045
 Normal Test Data Mean = 0.12066414207220078
 Normal Test Data Standard Deviation = 0.04088136553764343
 Abnormal Test Data Mean = 0.2277468740940094
 Abnormal Test Data Standard Deviation = 0.036684613674879074
 Precision = 0.9560315670800451
 Recall = 0.848
 0.9045
 0.12066414
 0.040881366
 0.22774687
 0.036684614
 0.9560315670800451
 0.848
 0.9045 0.12066414 0.040881366 0.22774687 0.036684614 0.9560315670800451 0.848

```
print("Threshold =", valid_data_best_threshold)
```

```
Threshold = 0.16348564155399803
```

```
print(confusion_matrix(test_labels_T_F, preds))
```

```
[[961  39]
 [152 848]]
```

Extra accuracy info

Just informative. Please record the above accuracy.

Accuracy on the test data with threshold set based on (threshold2 + threshold3) / 2 =
 Average of (mean + std of the distribution of the reconstruction losses of the normal
 validation data) and (mean - std of the distribution of the reconstruction losses of the
 abnormal validation data)

```
preds = predict(variational_ae, test_data, Avg_of_threshold_2_3)
print_stats(preds, test_labels_T_F)
```

Confusion Matrix:

		T
		1031 969
prediction: F	F	[[916 84]
	T	[115 885]] 1000

Accuracy = 0.9005

Normal Test Data Mean = 0.12066414207220078

Normal Test Data Standard Deviation = 0.04088136553764343

Abnormal Test Data Mean = 0.2277468740940094

Abnormal Test Data Standard Deviation = 0.036684613674879074

Precision = 0.913312693498452

Recall = 0.885

0.9005

0.12066414

0.040881366

0.22774687

0.036684614

0.913312693498452

0.885

0.9005 0.12066414 0.040881366 0.22774687 0.036684614 0.913312693498452 0.885

Accuracy on the test data with threshold set based on the mean of the training data MAE
 reconstruction losses + 2.5 std

```

preds = predict(variational_ae, test_data, threshold_train_mean_2_5_std)
print_stats(preds, test_labels_T_F)

Confusion Matrix:
 prediction: F      T
               556   1444
 label: F    [[535   465]    1000
              T    [21   979]]   1000
Accuracy = 0.757
Normal Test Data Mean = 0.12066414207220078
Normal Test Data Standard Deviation = 0.04088136553764343
Abnormal Test Data Mean = 0.2277468740940094
Abnormal Test Data Standard Deviation = 0.036684613674879074
Precision = 0.6779778393351801
Recall = 0.979
0.757
0.12066414
0.040881366
0.22774687
0.036684614
0.6779778393351801
0.979
0.757 0.12066414 0.040881366 0.22774687 0.036684614 0.6779778393351801 0.979

```

Extra Info

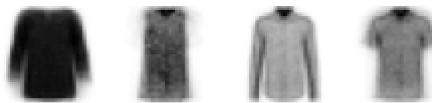
Giving the VAE codings (please see book) (Just informative, not the goal here)

```

def plot_multiple_images(images, n_cols=None):
    n_cols = n_cols or len(images)
    n_rows = (len(images) - 1) // n_cols + 1
    if images.shape[-1] == 1:
        images = np.squeeze(images, axis=-1)
    plt.figure(figsize=(n_cols, n_rows))
    for index, image in enumerate(images):
        plt.subplot(n_rows, n_cols, index + 1)
        plt.imshow(image, cmap="binary")
        plt.axis("off")

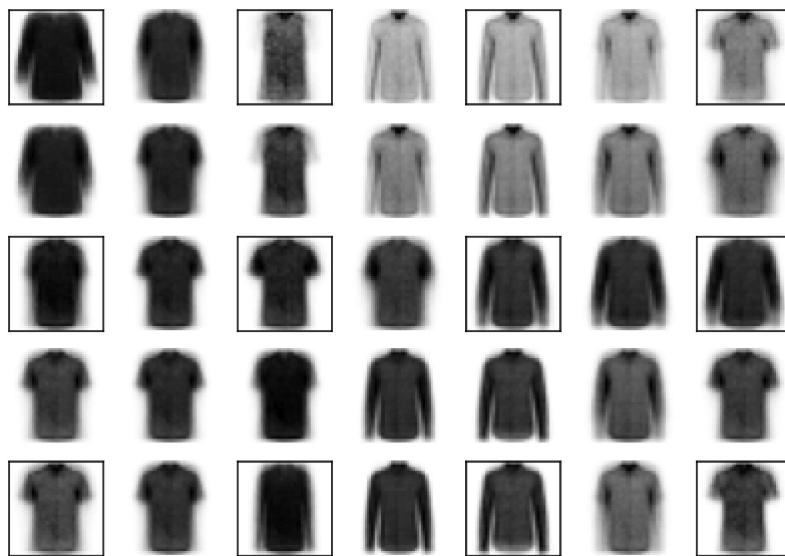
codings = tf.random.normal(shape=[12, codings_size])
images = variational_decoder(codings).numpy()
plot_multiple_images(images, 4)
# save_fig("vae_generated_images_plot", tight_layout=False)

```



```
codings_grid = tf.reshape(codings, [1, 3, 4, codings_size])
larger_grid = tf.image.resize(codings_grid, size=[5, 7])
interpolated_codings = tf.reshape(larger_grid, [-1, codings_size])
images = variational_decoder(interpolated_codings).numpy()

plt.figure(figsize=(7, 5))
for index, image in enumerate(images):
    plt.subplot(5, 7, index + 1)
    if index%7%2==0 and index//7%2==0:
        plt.gca().get_xaxis().set_visible(False)
        plt.gca().get_yaxis().set_visible(False)
    else:
        plt.axis("off")
    plt.imshow(image, cmap="binary")
# save_fig("semantic_interpolation_plot", tight_layout=False)
```



✓ 2s completed at 3:19 PM

