

References: Chapter 17 of Geron's book. For 1-Dim plots, Keras tutorial :

<https://www.tensorflow.org/tutorials/generative/autoencoder>

This file trains a VAE with the instances of the normal digits in the training data.

Then, it measures the reconstruction loss for the digits in the test data.

The reconstruction loss for the instances of the abnormal digits in the test data is higher.

A threshold is determined based on the distribution of the reconstruction losses of the normal training data (threshold = mean + 2.5*std of this distribution).

Then, if the reconstruction loss of a digit in the test data is higher than this threshold, it is classified as abnormal.

By comparing with the known labels of test data (with T for normal digit(s) and F for abnormal digit(s)), the confusion matrix and the accuracy is calculated.

```
import sklearn
import tensorflow as tf
from tensorflow import keras
import numpy as np

import matplotlib as mpl
import matplotlib.pyplot as plt
mpl.rc('axes', labelsiz=14)
mpl.rc('xtick', labelsiz=12)
mpl.rc('ytick', labelsiz=12)

from sklearn.metrics import accuracy_score, precision_score, recall_score, confusion_matrix
from sklearn.model_selection import train_test_split
```

Loading the MNIST data and forming arrays of the normal training data, the validation data (normal and abnormal) and the test data (normal and abnormal)

```
#Labels
# 0 T-shirt/top
# 1 Trouser
# 2 Pullover
# 3 Dress
# 4 Coat
# 5 Sandal
# 6 Shirt
# 7 Sneaker
# 8 Bag
# 9 Ankle boot

n11 = 6
n12 = 6
```

```
abn1 = 7
abn2 = 7
```

```
(x_train_0, y_train_0), (x_test, y_test) = keras.datasets.fashion_mnist.load_data()
```

```
x_train_0 = x_train_0.astype(np.float32) / 255
```

```
x_test = x_test.astype(np.float32) / 255
```

```
train_size = x_train_0.shape[0] * 9 // 10
```

```
x_train, x_valid, y_train, y_valid = train_test_split(x_train_0, y_train_0, train_size = train_size)
```

```
normal_data = x_train[(y_train == n11) | (y_train == n12)] # Normal training data (Normal digits
```

```
normal_labels = y_train[(y_train == n11) | (y_train == n12)]
```

```
valid_data = x_valid[(y_valid == abn1) | (y_valid == abn2) | (y_valid == n11) | (y_valid == n12)] #
```

```
valid_labels = y_valid[(y_valid == abn1) | (y_valid == abn2) | (y_valid == n11) | (y_valid == n12)]
```

```
test_data = x_test[(y_test == abn1) | (y_test == abn2) | (y_test == n11) | (y_test == n12)] # Test d
```

```
test_labels = y_test[(y_test == abn1) | (y_test == abn2) | (y_test == n11) | (y_test == n12)]
```

```
test_labels_T_F = np.where((test_labels == n11) | (test_labels == n12), True, False)
```

```
# Array of T and F, T where test digits are normal and F where test digits are abnormal
```

```
valid_labels_T_F = np.where((valid_labels == n11) | (valid_labels == n12), True, False)
```

```
# Array of T and F, T where test digits are normal and F where test digits are abnormal
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-labels-id:
```

```
32768/29515 [=====] - 0s 0us/step
```

```
40960/29515 [=====] - 0s 0us/step
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-images-id:
```

```
26427392/26421880 [=====] - 1s 0us/step
```

```
26435584/26421880 [=====] - 1s 0us/step
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-labels-idx:
```

```
16384/5148 [=====]
```

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-images-idx:
```

```
4423680/4422102 [=====] - 0s 0us/step
```

```
4431872/4422102 [=====] - 0s 0us/step
```

```
normal_data.shape, normal_labels.shape, valid_data.shape, valid_labels.shape, test_data.shape, test_label
```

```
((5383, 28, 28), (5383,), (1222, 28, 28), (1222,), (2000, 28, 28), (2000,))
```

```
normal_test_data = test_data[(test_labels == n11) | (test_labels == n12)] # The normal digit
```

```
abnormal_test_data = test_data[(test_labels == abn1) | (test_labels == abn2)] # The abnormal dig
```

```
normal_test_labels = test_labels[(test_labels == n11) | (test_labels == n12)] # Their labels
```

```
abnormal_test_labels = test_labels[(test_labels == abn1) | (test_labels == abn2)] # Their labels
```

```
normal_test_data.shape, abnormal_test_data.shape
```

```
((1000, 28, 28), (1000, 28, 28))
```

```
normal_valid_data = valid_data[(valid_labels == n11) | (valid_labels == n12)] # The normal d
```

```

abnormal_valid_data = valid_data[(valid_labels == abn1) | (valid_labels == abn2)] # The abnormal
normal_valid_labels = valid_labels[(valid_labels == n11) | (valid_labels == n12)] # Their labels
abnormal_valid_labels = valid_labels[(valid_labels == abn1) | (valid_labels == abn1)] # Their labels

normal_valid_data.shape, abnormal_valid_data.shape

((617, 28, 28), (605, 28, 28))

```

▼ Building and training the network

```

K = keras.backend
# def rounded_accuracy(y_true, y_pred):
#     return keras.metrics.binary_accuracy(tf.round(y_true), tf.round(y_pred))

# For details please see Geron's book. Uses the reparametrization trick to do stochastic
# sampling from the MVN distribution, while allowing the 2 parallel layers containing the
# means and stds of the MVN distribution for each dimension to be trained via
# backpropogation of the error signal.
class Sampling(keras.layers.Layer):
    def call(self, inputs):
        mean, log_var = inputs
        return K.random_normal(tf.shape(log_var)) * K.exp(log_var / 2) + mean

# For details please see Geron's book.
codings_size = 16 # The number of dimensions of the MVN distribution in the sampling layer

inputs = keras.layers.Input(shape=[28, 28])
z = keras.layers.Flatten()(inputs)
z = keras.layers.Dense(256, activation="selu")(z)
z = keras.layers.Dense(128, activation="selu")(z)
z = keras.layers.Dense(64, activation="selu")(z)

# Parallel layers at the end of the encoder for means
# and standard deviations of the Multivariate Normal (MVN) distribution
# in the dimensions of the coding size (here 32).
codings_mean = keras.layers.Dense(codings_size)(z)
codings_log_var = keras.layers.Dense(codings_size)(z)

# Sampling layer at the end of the encoder
codings = Sampling()([codings_mean, codings_log_var])
variational_encoder = keras.models.Model(
    inputs=[inputs], outputs=[codings_mean, codings_log_var, codings])

decoder_inputs = keras.layers.Input(shape=[codings_size])
x = keras.layers.Dense(64, activation="selu")(decoder_inputs)
x = keras.layers.Dense(128, activation="selu")(x)
x = keras.layers.Dense(256, activation="selu")(x)
x = keras.layers.Dense(28 * 28, activation="sigmoid")(x)
outputs = keras.layers.Reshape([28, 28])(x)
variational_decoder = keras.models.Model(inputs=[decoder_inputs], outputs=[outputs])

```

```
_, _, codings = variational_encoder(inputs)
reconstructions = variational_decoder(codings)
variational_ae = keras.models.Model(inputs=[inputs], outputs=[reconstructions])

# The latent loss function
# latent_loss = -0.5 * K.sum(
#     1 + codings_log_var - K.exp(codings_log_var) - K.square(codings_mean),
#     axis=-1)

# Add the latent loss to the reconstruction loss
# variational_ae.add_loss(K.mean(latent_loss) / 784.)

# For the reconstruction loss binary cross-entropy loss is used.
# For details please see Chapter 17 of Geron's book (Stacked AE and VAE sections)
variational_ae.compile(loss="binary_crossentropy", optimizer="rmsprop")

checkpoint_cb = keras.callbacks.ModelCheckpoint("wo_latent_VAE_model", monitor="val_loss", save_best_only=True)

history = variational_ae.fit(normal_data, normal_data, epochs=100, batch_size=128, callbacks=[checkpoint_cb],
                             validation_data=(normal_valid_data, normal_valid_data), shuffle=True)
```

```
Epoch 72/100
43/43 [=====] - 0s 9ms/step - loss: 0.3529 - val_loss: 0.3633
Epoch 73/100
43/43 [=====] - 0s 9ms/step - loss: 0.3533 - val_loss: 0.3642
Epoch 74/100
43/43 [=====] - 0s 9ms/step - loss: 0.3527 - val_loss: 0.3609
Epoch 75/100
43/43 [=====] - 0s 8ms/step - loss: 0.3531 - val_loss: 0.3647
Epoch 76/100
43/43 [=====] - 0s 9ms/step - loss: 0.3526 - val_loss: 0.3658
Epoch 77/100
43/43 [=====] - 0s 9ms/step - loss: 0.3527 - val_loss: 0.3654
Epoch 78/100
43/43 [=====] - 0s 8ms/step - loss: 0.3526 - val_loss: 0.3642
Epoch 79/100
43/43 [=====] - 0s 9ms/step - loss: 0.3523 - val_loss: 0.3639
Epoch 80/100
43/43 [=====] - 0s 9ms/step - loss: 0.3524 - val_loss: 0.3647
Epoch 81/100
43/43 [=====] - 0s 9ms/step - loss: 0.3522 - val_loss: 0.3638
Epoch 82/100
40/43 [=====>...] - ETA: 0s - loss: 0.3523INFO:tensorflow:Assets written to disk
43/43 [=====] - 3s 65ms/step - loss: 0.3523 - val_loss: 0.3591
Epoch 83/100
43/43 [=====] - 0s 9ms/step - loss: 0.3522 - val_loss: 0.3634
Epoch 84/100
43/43 [=====] - 0s 9ms/step - loss: 0.3518 - val_loss: 0.3627
Epoch 85/100
43/43 [=====] - 0s 8ms/step - loss: 0.3519 - val_loss: 0.3613
Epoch 86/100
43/43 [=====] - 0s 9ms/step - loss: 0.3519 - val_loss: 0.3656
Epoch 87/100
43/43 [=====] - 0s 9ms/step - loss: 0.3518 - val_loss: 0.3615
Epoch 88/100
43/43 [=====] - 0s 9ms/step - loss: 0.3515 - val_loss: 0.3600
Epoch 89/100
43/43 [=====] - 0s 9ms/step - loss: 0.3516 - val_loss: 0.3603
```

```

Epoch 90/100
43/43 [=====] - 0s 8ms/step - loss: 0.3514 - val_loss: 0.3663
Epoch 91/100
43/43 [=====] - 0s 9ms/step - loss: 0.3515 - val_loss: 0.3644
Epoch 92/100
43/43 [=====] - ETA: 0s - loss: 0.3513INFO:tensorflow:Assets written t
43/43 [=====] - 3s 71ms/step - loss: 0.3513 - val_loss: 0.3586
Epoch 93/100
43/43 [=====] - 0s 8ms/step - loss: 0.3512 - val_loss: 0.3616
Epoch 94/100
43/43 [=====] - 0s 8ms/step - loss: 0.3512 - val_loss: 0.3601
Epoch 95/100
43/43 [=====] - 0s 8ms/step - loss: 0.3507 - val_loss: 0.3669
Epoch 96/100
43/43 [=====] - 0s 9ms/step - loss: 0.3512 - val_loss: 0.3646
Epoch 97/100
43/43 [=====] - 0s 8ms/step - loss: 0.3507 - val_loss: 0.3623
Epoch 98/100
43/43 [=====] - 0s 8ms/step - loss: 0.3510 - val_loss: 0.3613
Epoch 99/100
43/43 [=====] - 0s 8ms/step - loss: 0.3508 - val_loss: 0.3604


```

```

plt.plot(history.history["loss"], label="Training Loss")
plt.plot(history.history["val_loss"], label="Validation Loss")
plt.legend()

```

<matplotlib.legend.Legend at 0x7f4e2abcd510>



```

model = variational_ae
model.summary(expand_nested=True, show_trainable=True)

```

Model: "model_2"

Layer (type)	Output Shape	Param #	Trainable
input_1 (InputLayer)	[(None, 28, 28)]	0	Y
model (Functional)	[(None, 16), (None, 16), (None, 16)]	244192	Y
input_1 (InputLayer)	[(None, 28, 28)]	0	Y

flatten (Flatten)	(None, 784)	0	Y
dense (Dense)	(None, 256)	200960	Y
dense_1 (Dense)	(None, 128)	32896	Y
dense_2 (Dense)	(None, 64)	8256	Y
dense_3 (Dense)	(None, 16)	1040	Y
dense_4 (Dense)	(None, 16)	1040	Y
sampling (Sampling)	(None, 16)	0	Y

model_1 (Functional)	(None, 28, 28)	243920	Y

input_2 (InputLayer)	[(None, 16)]	0	Y
dense_5 (Dense)	(None, 64)	1088	Y
dense_6 (Dense)	(None, 128)	8320	Y
dense_7 (Dense)	(None, 256)	33024	Y
dense_8 (Dense)	(None, 784)	201488	Y
reshape (Reshape)	(None, 28, 28)	0	Y

=====			
Total params: 488,112			
Trainable params: 488,112			
Non-trainable params: 0			
=====			

```
model_encoder = variational_encoder
# model_encoder.summary(expand_nested=True, show_trainable=True)
```

```
model_decoder = variational_decoder
# model_decoder.summary(expand_nested=True, show_trainable=True)
```

```
model_layers = np.array(model.layers)
n_layers = model_layers.shape[0]
# np.concatenate((np.arange(n_layers).reshape(n_layers,1), model_layers.reshape(n_layers,1)), axis = 1)
```

The original and reconstructed images for the first 30 instances of the normal training data, validation data, normal validation data, abnormal validation data, test data, normal test data, and abnormal test data

```
def plot_image(image):
    plt.imshow(image, cmap="binary")
    plt.axis("off")
```

```
def show_reconstructions(model, images, n_images=5):
    reconstructions = model.predict(images[:n_images])
    fig = plt.figure(figsize=(n_images * 1.5, 3))
    for image_index in range(n_images):
        plt.subplot(2, n_images, 1 + image_index)
        plot_image(images[image_index])
        plt.subplot(2, n_images, 1 + n_images + image_index)
        plot_image(reconstructions[image_index])
```

```
show_reconstructions(variational_ae, normal_data, 30)
plt.show()
```



```
show_reconstructions(variational_ae, valid_data, 30)
plt.show()
```



```
show_reconstructions(variational_ae, normal_valid_data, 30)
plt.show()
```



```
show_reconstructions(variational_ae, abnormal_valid_data, 30)
plt.show()
```



```
show_reconstructions(variational_ae, test_data, 30)
plt.show()
```



```
show_reconstructions(variational_ae, normal_test_data, 30)
plt.show()
```




```
show_reconstructions(variational_ae, abnormal_test_data, 30)
plt.show()
```



1-Dim plot of pixels of the first normal test data

```
reconstructions_n1_test = variational_ae.predict(normal_test_data)
```


```
plt.figure(figsize=(25,7))
plt.plot(normal_test_data[0].ravel(), 'r')
plt.plot(reconstructions_n1_test[0].ravel(), 'g')
plt.fill_between(np.arange(28*28), reconstructions_n1_test[0].ravel(), normal_test_data[0].ravel(), color='b')
plt.legend(labels=["Input", "Reconstruction", "Error"])
plt.show()
```

1-Dim plot of pixels of the first abnormal test data

```
reconstructions_abn_test = variational_ae.predict(abnormal_test_data)
```

```
plt.figure(figsize=(25,7))
plt.plot(abnormal_test_data[0].ravel(), 'r')
plt.plot(reconstructions_abn_test[0].ravel(), 'g')
plt.fill_between(np.arange(28*28), reconstructions_abn_test[0].ravel(), normal_test_data[0].ravel(), color='b')
plt.legend(labels=["Input", "Reconstruction", "Error"])
plt.show()
```



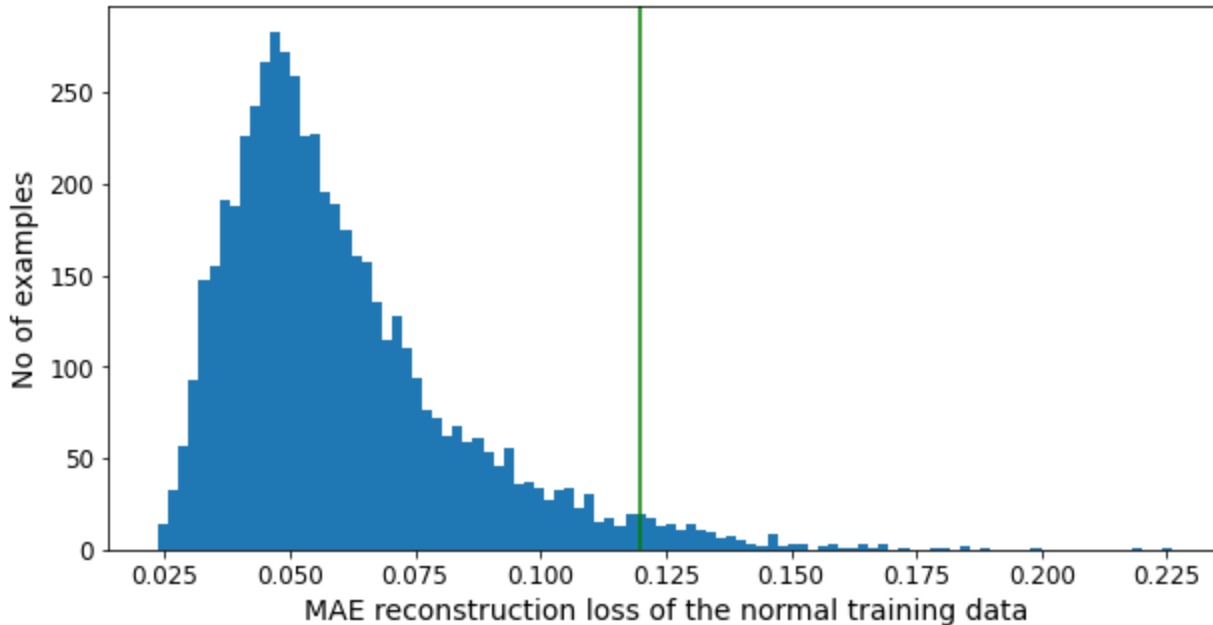
► Distributions of the reconstruction losses and the calculation of the threshold.

Distribution of the reconstruction losses of the normal training data

```

reconstructions = variational_ae.predict(normal_data)
train_loss = tf.keras.losses.mae(reconstructions.reshape(-1, 784), normal_data.reshape(-1, 784))
plt.figure(figsize=(10,5))
plt.hist(train_loss[None,:], bins=100)
threshold1 = np.mean(train_loss) + 2.5*np.std(train_loss)
plt.axvline(threshold1,c='g')
plt.xlabel("MAE reconstruction loss of the normal training data")
plt.ylabel("No of examples")
plt.show()

```



```

print("Mean: ", np.mean(train_loss))
print("Std: ", np.std(train_loss))

```

```

Mean:  0.060375232
Std:   0.023860663

```

```

threshold_train_mean_2_5_std = np.mean(train_loss) + 2.5*np.std(train_loss)
print("Threshold based on the mean of the training data MAE reconstruction losses + 2.5 std: ", threshold_train_mean_2_5_std)

```

```

Threshold based on the mean of the training data MAE reconstruction losses + 2.5 std:  0.120026890

```

```

threshold1 = threshold_train_mean_2_5_std

```

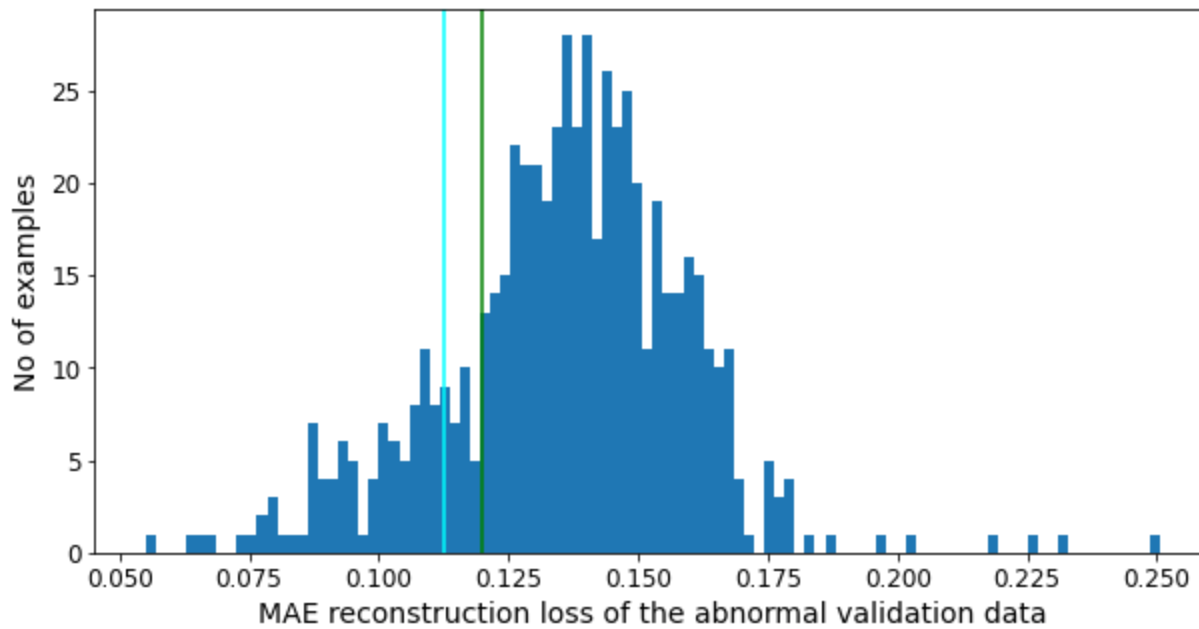
Distribution of the reconstruction losses of the abnormal validation data

```

reconstructions = variational_ae.predict(abnormal_valid_data)
abn_valid_loss = tf.keras.losses.mae(reconstructions.reshape(-1,784), abnormal_valid_data.reshape(-1,784))
plt.figure(figsize=(10,5))
plt.hist(abn_valid_loss[None, :], bins=100)
threshold2 = np.mean(abn_valid_loss) - np.std(abn_valid_loss)
plt.axvline(threshold2,c='cyan')
plt.axvline(threshold1,c='g')

```

```
plt.xlabel("MAE reconstruction loss of the abnormal validation data")
plt.ylabel("No of examples")
plt.show()
```



```
abnormal_valid_mean_loss = np.mean(abn_valid_loss)
```

```
abnormal_valid_mean_loss , np.std(abn_valid_loss)
```


```
(0.13596502, 0.023510419)
```

```
threshold2 = abnormal_valid_mean_loss - np.std(abn_valid_loss)
print("Threshold2: ", threshold2)
```

```
Threshold2: 0.1124546
```

Distribution of the reconstruction losses of the normal validation data

```
reconstructions = variational_ae.predict(normal_valid_data)
nl_valid_loss = tf.keras.losses.mae(reconstructions.reshape(-1,784), normal_valid_data.reshape(-1,784))
plt.figure(figsize=(10,5))
plt.hist(nl_valid_loss[None, :], bins=100)
threshold3 = np.mean(nl_valid_loss) + np.std(nl_valid_loss)
plt.axvline(threshold3, c='magenta')
plt.axvline(threshold2, c='cyan')
plt.axvline(threshold1, c='g')
plt.xlabel("MAE reconstruction loss of the normal validation data")
plt.ylabel("No of examples")
plt.show()
```



```
normal_valid_mean_loss = np.mean(nl_valid_loss)

normal_valid_mean_loss , np.std(nl_valid_loss)

(0.06431188, 0.026296)
```

```
threshold3 = normal_valid_mean_loss + np.std(nl_valid_loss)
print("Threshold3: ", threshold3)
```

```
Threshold3:  0.09060788
```

Calculation of a preliminary threshold based on $(\text{threshold2} + \text{threshold3}) / 2 = \text{Average of (mean + std of the distribution of the reconstruction losses of the normal validation data) and (mean - std of the distribution of the reconstruction losses of the abnormal validation data)}$

```
Avg_of_threshold_2_3 = (threshold2 + threshold3)/2
print("Average of threshold 2 and 3: ", Avg_of_threshold_2_3)
```

```
Average of threshold 2 and 3:  0.10153123736381531
```

```
threshold4 = Avg_of_threshold_2_3
```

Calculation of the threshold that gives the best accuracy on the validation data and set this as the threshold.

```
def predict(model, data, threshold):
    reconstructions = model.predict(data)
    loss = tf.keras.losses.mae(reconstructions.reshape(-1, 784), data.reshape(-1, 784))
    return tf.math.less(loss, threshold)
```

```
increment = (abnormal_valid_mean_loss- normal_valid_mean_loss)/100
thresholds = np.arange(normal_valid_mean_loss, abnormal_valid_mean_loss, increment)
thrs_size = thresholds.shape[0]
```

```

accuracies = np.zeros(thrs_size)
for i in range(thrs_size):
    preds = predict(variational_ae, valid_data, thresholds[i])
    accuracies[i] = accuracy_score(preds, valid_labels_T_F)
argmax = np.argmax(accuracies)
valid_data_best_threshold = thresholds[argmax]
print("The best threshold based on validation data: ", valid_data_best_threshold)

```

The best threshold based on validation data: 0.09225660249590893

```

thr_acc = np.zeros((thrs_size, 2))
thr_acc[:, 0] = thresholds
thr_acc[:, 1] = accuracies
thr_acc[argmax-2:argmax+3]

array([[0.09082354, 0.91080196],
       [0.09154007, 0.91243863],
       [0.0922566 , 0.91325696],
       [0.09297313, 0.91243863],
       [0.09368967, 0.91162029]])

```

threshold5 = valid_data_best_threshold

threshold = threshold5

▼ Distribution of the reconstruction losses of all the validation data (normal and abnormal)


The blue line is threshold4 (= the average of threshold3 [magenta] and threshold2 [cyan]).

The red line is the threshold that gives the best accuracy for the validation data.

```

reconstructions = variational_ae.predict(valid_data)
valid_loss = tf.keras.losses.mae(reconstructions.reshape(-1,784), valid_data.reshape(-1,784))
plt.figure(figsize=(20,5))
plt.hist(valid_loss[None, :], bins=100)
plt.axvline(threshold, c='r')
plt.axvline(threshold4, c='b')
plt.axvline(threshold2, c='cyan')
plt.axvline(threshold3, c='magenta')
plt.axvline(threshold1, c='green')
plt.xlabel("MAE reconstruction loss of the validation data")
plt.ylabel("No of examples")
plt.show()

```




▼ Distribution of the reconstruction losses of the test data (normal and abnormal)

The blue line is threshold4 (= the average of threshold3 [magenta] and threshold2 [cyan]).

The red line is the threshold that gives the best accuracy for the validation data.

```
reconstructions = variational_ae.predict(test_data)
test_loss = tf.keras.losses.mae(reconstructions.reshape(-1,784), test_data.reshape(-1,784))
plt.figure(figsize=(20,5))
plt.hist(test_loss[None, :], bins=100)
plt.axvline(threshold, c='r')
plt.axvline(threshold4, c='b')
plt.axvline(threshold2, c='cyan')
plt.axvline(threshold3, c='magenta')
plt.axvline(threshold1, c='green')
plt.xlabel("MAE reconstruction loss of the test data")
plt.ylabel("No of examples")
plt.show()
```



▼ Mean and standard deviation of reconstruction losses for normal and abnormal test data

```
reconstructions = variational_ae.predict(normal_test_data)
```

```

nl_test_loss = tf.keras.losses.mae(reconstructions.reshape(-1,784), normal_test_data.reshape(-1,784))
np.mean(nl_test_loss) , np.std(nl_test_loss)

(0.06429024, 0.02582075)

reconstructions = variational_ae.predict(abnormal_test_data)
abn_test_loss = tf.keras.losses.mae(reconstructions.reshape(-1,784), abnormal_test_data.reshape(-1,784))
np.mean(abn_test_loss) , np.std(abn_test_loss)

(0.13432911, 0.02094467)

```

Calculation of the accuracy and the confusion matrix on the test data with threshold set based on the best threshold from the validation data

```

# def predict(model, data, threshold):
#     reconstructions = model.predict(data)
#     loss = tf.keras.losses.mae(reconstructions.reshape(-1, 784), data.reshape(-1, 784))
#     return tf.math.less(loss, threshold)

def print_stats(predictions, labels):
    cf = confusion_matrix(labels, predictions)
    print("Confusion Matrix: \n prediction: F      T ")
    print("      {}   {}".format(preds[preds == False].shape[0], preds[preds == True].shape[0]))
    print("label: F   [[{}   {}]]   {}".format(cf[0,0], cf[0,1], test_labels_T_F[test_labels_T_F == False].shape[0]))
    print("      T   [[{}   {}]]   {}".format(cf[1,0], cf[1,1], test_labels_T_F[test_labels_T_F == True].shape[0]))
    print("Accuracy = {}".format(accuracy_score(labels, predictions)))
    print("Normal Test Data Mean = {}".format(np.mean(nl_test_loss)))
    print("Normal Test Data Standard Deviation = {}".format(np.std(nl_test_loss)))
    print("Abnormal Test Data Mean = {}".format(np.mean(abn_test_loss)))
    print("Abnormal Test Data Standard Deviation = {}".format(np.std(abn_test_loss)))
    print("Precision = {}".format(precision_score(labels, predictions)))
    print("Recall = {}".format(recall_score(labels, predictions)))
    print(accuracy_score(labels, predictions))
    print(np.mean(nl_test_loss))
    print(np.std(nl_test_loss))
    print(np.mean(abn_test_loss))
    print(np.std(abn_test_loss))
    print(precision_score(labels, predictions))
    print(recall_score(labels, predictions))
    print(accuracy_score(labels, predictions), np.mean(nl_test_loss), np.std(nl_test_loss), np.mean(abn_test_loss), np.std(abn_test_loss), precision_score(labels, predictions), recall_score(labels, predictions))

```

```

preds = predict(variational_ae, test_data, threshold)
print_stats(preds, test_labels_T_F)

```

```

☞ Confusion Matrix:
prediction: F      T
           1085   915
label: F   [[959   41]]   1000
           T   [126   874]]   1000
Accuracy = 0.9165
Normal Test Data Mean = 0.06429024040699005
Normal Test Data Standard Deviation = 0.02582075074315071

```

```

Abnormal Test Data Mean = 0.13432911038398743
Abnormal Test Data Standard Deviation = 0.020944669842720032
Precision = 0.9551912568306011
Recall = 0.874
0.9165
0.06429024
0.02582075
0.13432911
0.02094467
0.9551912568306011
0.874
0.9165 0.06429024 0.02582075 0.13432911 0.02094467 0.9551912568306011 0.874

```

```
print("Threshold =", valid_data_best_threshold)
```

```
Threshold = 0.09225660249590893
```

```
print(confusion_matrix(test_labels_T_F, preds))
```

```

[[959  41]
 [126 874]]

```

Extra accuracy info

Just informative. Please record the above accuracy.

Accuracy on the test data with threshold set based on $(\text{threshold}_2 + \text{threshold}_3) / 2$ = Average of
 ▼ (mean + std of the distribution of the reconstruction losses of the normal validation data) and
 (mean - std of the distribution of the reconstruction losses of the abnormal validation data)

```

preds = predict(variational_ae, test_data, Avg_of_threshold_2_3)
print_stats(preds, test_labels_T_F)

```

```

Confusion Matrix:
prediction: F      T
           1010   990
label: F    [[925   75]   1000
           T    [ 85  915]]   1000

```

```

Accuracy = 0.92
Normal Test Data Mean = 0.06429024040699005
Normal Test Data Standard Deviation = 0.02582075074315071
Abnormal Test Data Mean = 0.13432911038398743
Abnormal Test Data Standard Deviation = 0.020944669842720032
Precision = 0.9242424242424242
Recall = 0.915
0.92
0.06429024
0.02582075
0.13432911
0.02094467
0.9242424242424242
0.915
0.92 0.06429024 0.02582075 0.13432911 0.02094467 0.9242424242424242 0.915

```


Accuracy on the test data with threshold set based on the mean of the training data MAE reconstruction losses + 2.5 std

```
preds = predict(variational_ae, test_data, threshold_train_mean_2_5_std)
print_stats(preds, test_labels_T_F)
```

Confusion Matrix:

prediction:	F	T	
	821	1179	
label: F	[[780	220]	1000
T	[41	959]]	1000

Accuracy = 0.8695

Normal Test Data Mean = 0.06429024040699005

Normal Test Data Standard Deviation = 0.02582075074315071

Abnormal Test Data Mean = 0.13432911038398743

Abnormal Test Data Standard Deviation = 0.020944669842720032

Precision = 0.813401187446989

Recall = 0.959

0.8695

0.06429024

0.02582075

0.13432911

0.02094467

0.813401187446989

0.959


0.8695 0.06429024 0.02582075 0.13432911 0.02094467 0.813401187446989 0.959

Extra Info

Giving the VAE codings (please see book) (Just informative, not the goal here)


```
def plot_multiple_images(images, n_cols=None):
    n_cols = n_cols or len(images)
    n_rows = (len(images) - 1) // n_cols + 1
    if images.shape[-1] == 1:
        images = np.squeeze(images, axis=-1)
    plt.figure(figsize=(n_cols, n_rows))
    for index, image in enumerate(images):
        plt.subplot(n_rows, n_cols, index + 1)
        plt.imshow(image, cmap="binary")
        plt.axis("off")
```

```
codings = tf.random.normal(shape=[12, codings_size])
images = variational_decoder(codings).numpy()
plot_multiple_images(images, 4)
# save_fig("vae_generated_images_plot", tight_layout=False)
```



```
codings_grid = tf.reshape(codings, [1, 3, 4, codings_size])
larger_grid = tf.image.resize(codings_grid, size=[5, 7])
interpolated_codings = tf.reshape(larger_grid, [-1, codings_size])
images = variational_decoder(interpolated_codings).numpy()
```

```
plt.figure(figsize=(7, 5))
for index, image in enumerate(images):
    plt.subplot(5, 7, index + 1)
    if index%7%2==0 and index//7%2==0:
        plt.gca().get_xaxis().set_visible(False)
        plt.gca().get_yaxis().set_visible(False)
    else:
        plt.axis("off")
    plt.imshow(image, cmap="binary")
# save_fig("semantic_interpolation_plot", tight_layout=False)
```



✓ 1s completed at 1:29 PM

