

Reference: Keras tutorial : <https://www.tensorflow.org/tutorials/generative/autoencoder> . Chapter 17 of Geron's book.

This file trains an autoencoder with the instances of the normal digit in the training data.

Then, it measures the reconstruction loss for the digits in the test data.

The reconstruction loss for the instances of the abnormal digit in the test data is higher.

A threshold is determined based on the distribution of the reconstruction losses of the normal training data (threshold = mean + 2.5*std of this distribution).

Then, if the reconstruction loss of a digit in the test data is higher than this threshold, it is classified as abnormal.

By comparing with the known labels of test data (with T for normal digit(s) and F for abnormal digit(s)), the confusion matrix and the accuracy is calculated.

```
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import tensorflow as tf
from tensorflow import keras

from sklearn.metrics import accuracy_score, precision_score, recall_score, confusion_matrix
from sklearn.model_selection import train_test_split
from keras import layers, losses
from keras.datasets import mnist
from keras.models import Model
# import cv2
```

Loading the MNIST data and forming arrays of the normal training data, validation data (normal and abnormal), and the test data (normal and abnormal)

```
#Labels
# 0 T-shirt/top
# 1 Trouser
# 2 Pullover
# 3 Dress
# 4 Coat
# 5 Sandal
# 6 Shirt
# 7 Sneaker
# 8 Bag
# 9 Ankle boot
```

```

nl1 = 9
nl2 = 9
abn = 3

(x_train_0, y_train_0), (x_test, y_test) = keras.datasets.fashion_mnist.load_data()

x_train_0 = x_train_0.astype(np.float32) / 255
x_test = x_test.astype(np.float32) / 255

train_size = x_train_0.shape[0] * 9 // 10

x_train, x_valid, y_train, y_valid = train_test_split(x_train_0, y_train_0, train_size = train_size)

normal_data = x_train[(y_train == nl1) | (y_train == nl2)]      # Normal training data (Normal)
normal_labels = y_train[(y_train == nl1) | (y_train == nl2)]

valid_data = x_valid[(y_valid == abn) | (y_valid == nl1) | (y_valid == nl2)]      # Validation
valid_labels = y_valid[(y_valid == abn) | (y_valid == nl1) | (y_valid == nl2)]

test_data = x_test[(y_test == abn) | (y_test == nl1) | (y_test == nl2)]      # Test data (both normal and abnormal)
test_labels = y_test[(y_test == abn) | (y_test == nl1) | (y_test == nl2)]

test_labels_T_F = np.where((test_labels == nl1) | (test_labels == nl2), True, False)
# Array of T and F, T where test digits are normal and F where test digits are abnormal

valid_labels_T_F = np.where((valid_labels == nl1) | (valid_labels == nl2), True, False)
# Array of T and F, T where test digits are normal and F where test digits are abnormal

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-32768/29515 [=====] - 0s 1us/step
40960/29515 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/train-26427392/26421880 [=====] - 0s 0us/step
26435584/26421880 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-16384/5148 [=====]
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/t10k-4423680/4422102 [=====] - 0s 0us/step
4431872/4422102 [=====] - 0s 0us/step

```



```

normal_data.shape, normal_labels.shape, valid_data.shape, valid_labels.shape, test_data.shape
((5392, 28, 28), (5392,), (1205, 28, 28), (1205,), (2000, 28, 28), (2000,))

normal_test_data = test_data[(test_labels == nl1) | (test_labels == nl2)]      # The normal digits
abnormal_test_data = test_data[test_labels == abn]                            # The abnormal digits
normal_test_labels = test_labels[(test_labels == nl1) | (test_labels == nl2)]    # Their label
abnormal_test_labels = test_labels[test_labels == abn]                         # Their label

```

```
normal_test_data.shape, abnormal_test_data.shape
```

```
((1000, 28, 28), (1000, 28, 28))
```

```
normal_valid_data = valid_data[(valid_labels == nl1) | (valid_labels == nl2)] # The normal  
abnormal_valid_data = valid_data[valid_labels == abn] # The abnormal  
normal_valid_labels = valid_labels[(valid_labels == nl1) | (valid_labels == nl2)] # Their labels  
abnormal_valid_labels = valid_labels[valid_labels == abn] # Their labels
```

```
normal_valid_data.shape, abnormal_valid_data.shape
```

```
((608, 28, 28), (597, 28, 28))
```

▼ Building and training the network

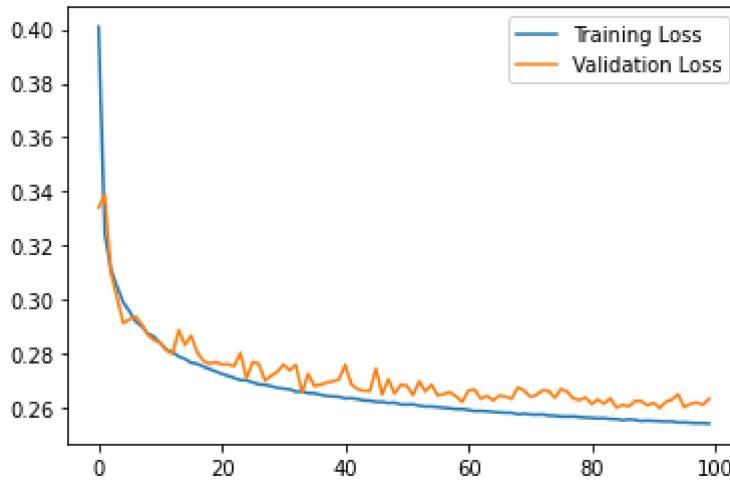
```
class AnomalyDetector(Model):  
    def __init__(self):  
        super(AnomalyDetector, self).__init__()  
        self.encoder = tf.keras.Sequential([  
            layers.Flatten(),  
            layers.Dense(256, activation="selu"),  
            layers.Dense(128, activation="selu"),  
            layers.Dense(64, activation="selu"),  
            layers.Dense(16, activation="selu")])  
  
        self.decoder = tf.keras.Sequential([  
            layers.Dense(64, activation="selu"),  
            layers.Dense(128, activation="selu"),  
            layers.Dense(256, activation="selu"),  
            layers.Dense(28*28, activation="sigmoid"),  
            layers.Reshape((28, 28))])  
  
    def call(self, x):  
        encoded = self.encoder(x)  
        decoded = self.decoder(encoded)  
        return decoded  
  
autoencoder = AnomalyDetector()  
  
# autoencoder.compile(optimizer='adam', loss='mae')  
autoencoder.compile(optimizer='rmsprop', loss='binary_crossentropy')  
  
checkpoint_cb = keras.callbacks.ModelCheckpoint("AE_model", monitor="val_loss", save_best_only=True)  
history = autoencoder.fit(normal_data, normal_data,
```

```
epochs=100,  
batch_size=128,  
validation_data=(normal_valid_data, normal_valid_data),  
callbacks=[checkpoint_cb],  
shuffle=True)  
  
Epoch 73/100  
43/43 [=====] - 0s 5ms/step - loss: 0.2576 - val_loss: 0.266  
Epoch 74/100  
43/43 [=====] - 0s 6ms/step - loss: 0.2572 - val_loss: 0.266  
Epoch 75/100  
43/43 [=====] - 0s 5ms/step - loss: 0.2571 - val_loss: 0.263  
Epoch 76/100  
43/43 [=====] - 0s 5ms/step - loss: 0.2568 - val_loss: 0.267  
Epoch 77/100  
43/43 [=====] - 0s 5ms/step - loss: 0.2568 - val_loss: 0.266  
Epoch 78/100  
43/43 [=====] - 0s 5ms/step - loss: 0.2568 - val_loss: 0.263  
Epoch 79/100  
43/43 [=====] - 0s 6ms/step - loss: 0.2566 - val_loss: 0.263  
Epoch 80/100  
43/43 [=====] - 0s 5ms/step - loss: 0.2564 - val_loss: 0.263  
Epoch 81/100  
35/43 [=====>.....] - ETA: 0s - loss: 0.2559INFO:tensorflow:Assets  
43/43 [=====] - 2s 42ms/step - loss: 0.2563 - val_loss: 0.263  
Epoch 82/100  
43/43 [=====] - 0s 5ms/step - loss: 0.2561 - val_loss: 0.263  
Epoch 83/100  
43/43 [=====] - 0s 5ms/step - loss: 0.2562 - val_loss: 0.261  
Epoch 84/100  
43/43 [=====] - 0s 5ms/step - loss: 0.2559 - val_loss: 0.263  
Epoch 85/100  
35/43 [=====>.....] - ETA: 0s - loss: 0.2560INFO:tensorflow:Assets  
43/43 [=====] - 2s 36ms/step - loss: 0.2559 - val_loss: 0.261  
Epoch 86/100  
43/43 [=====] - 0s 5ms/step - loss: 0.2554 - val_loss: 0.261  
Epoch 87/100  
43/43 [=====] - 0s 5ms/step - loss: 0.2557 - val_loss: 0.260  
Epoch 88/100  
43/43 [=====] - 0s 6ms/step - loss: 0.2555 - val_loss: 0.262  
Epoch 89/100  
43/43 [=====] - 0s 5ms/step - loss: 0.2551 - val_loss: 0.262  
Epoch 90/100  
43/43 [=====] - 0s 5ms/step - loss: 0.2553 - val_loss: 0.260  
Epoch 91/100  
43/43 [=====] - 0s 5ms/step - loss: 0.2552 - val_loss: 0.261  
Epoch 92/100  
35/43 [=====>.....] - ETA: 0s - loss: 0.2548INFO:tensorflow:Assets  
43/43 [=====] - 2s 36ms/step - loss: 0.2551 - val_loss: 0.261  
Epoch 93/100  
43/43 [=====] - 0s 5ms/step - loss: 0.2549 - val_loss: 0.262  
Epoch 94/100  
43/43 [=====] - 0s 5ms/step - loss: 0.2550 - val_loss: 0.263  
Epoch 95/100  
43/43 [=====] - 0s 5ms/step - loss: 0.2546 - val_loss: 0.265  
Epoch 96/100  
43/43 [=====] - 0s 5ms/step - loss: 0.2546 - val_loss: 0.265
```

```
43/43 [=====] - 0s 5ms/step - loss: 0.2544 - val_loss: 0.261
Epoch 97/100
43/43 [=====] - 0s 5ms/step - loss: 0.2543 - val_loss: 0.262
Epoch 98/100
43/43 [=====] - 0s 5ms/step - loss: 0.2543 - val_loss: 0.261
Epoch 99/100
43/43 [=====] - 0s 5ms/step - loss: 0.2543 - val_loss: 0.261
```

```
plt.plot(history.history["loss"], label="Training Loss")
plt.plot(history.history["val_loss"], label="Validation Loss")
plt.legend()
```

<matplotlib.legend.Legend at 0x7fb1e8607dd0>



```
model = autoencoder
model.summary(expand_nested=True, show_trainable=True)
```

Model: "anomaly_detector"

Layer (type)	Output Shape	Param #	Trainable
<hr/>			
sequential (Sequential)	(None, 16)	243152	Y

flatten (Flatten)	(None, 784)	0	Y
dense (Dense)	(None, 256)	200960	Y
dense_1 (Dense)	(None, 128)	32896	Y
dense_2 (Dense)	(None, 64)	8256	Y
dense_3 (Dense)	(None, 16)	1040	Y

sequential_1 (Sequential)	(None, 28, 28)	243920	Y

dense_4 (Dense)	(None, 64)	1088	Y
dense_5 (Dense)	(None, 128)	8320	Y

dense_6 (Dense)	(None, 256)	33024	Y	
dense_7 (Dense)	(None, 784)	201488	Y	
reshape (Reshape)	(None, 28, 28)	0	Y	

```
=====
Total params: 487,072
Trainable params: 487,072
Non-trainable params: 0
```

```
model_encoder = autoencoder.encoder
# model_encoder.summary(expand_nested=True, show_trainable=True)

model_decoder = autoencoder.decoder
# model_decoder.summary(expand_nested=True, show_trainable=True)

model_layers = np.array(model.layers)
n_layers = model_layers.shape[0]
# np.concatenate((np.arange(n_layers).reshape(n_layers,1), model_layers.reshape(n_layers,1)),
```

- The original and reconstructed images for the first 30 instances of the normal**
- ▼ **training data, validation data, normal validation data, abnormal validation data, test data, normal test data, and abnormal test data**

```
def plot_image(image):
    plt.imshow(image, cmap="binary")
    plt.axis("off")

def show_reconstructions(autoencoder, images, n_images=5):
    encoded_data = autoencoder.encoder(images[:n_images]).numpy()
    decoded_data = autoencoder.decoder(encoded_data).numpy()
    reconstructions = decoded_data
    fig = plt.figure(figsize=(n_images * 1.5, 3))
    for image_index in range(n_images):
        plt.subplot(2, n_images, 1 + image_index)
        plot_image(images[image_index])
        plt.subplot(2, n_images, 1 + n_images + image_index)
        plot_image(reconstructions[image_index])

show_reconstructions(autoencoder, normal_data, 30)
plt.show()
```



```
show_reconstructions(autoencoder, valid_data, 30)  
plt.show()
```



```
show_reconstructions(autoencoder, normal_valid_data, 30)  
plt.show()
```



```
show_reconstructions(autoencoder, abnormal_valid_data, 30)  
plt.show()
```



```
show_reconstructions(autoencoder, test_data, 30)  
plt.show()
```



```
show_reconstructions(autoencoder, normal_test_data, 30)  
plt.show()
```



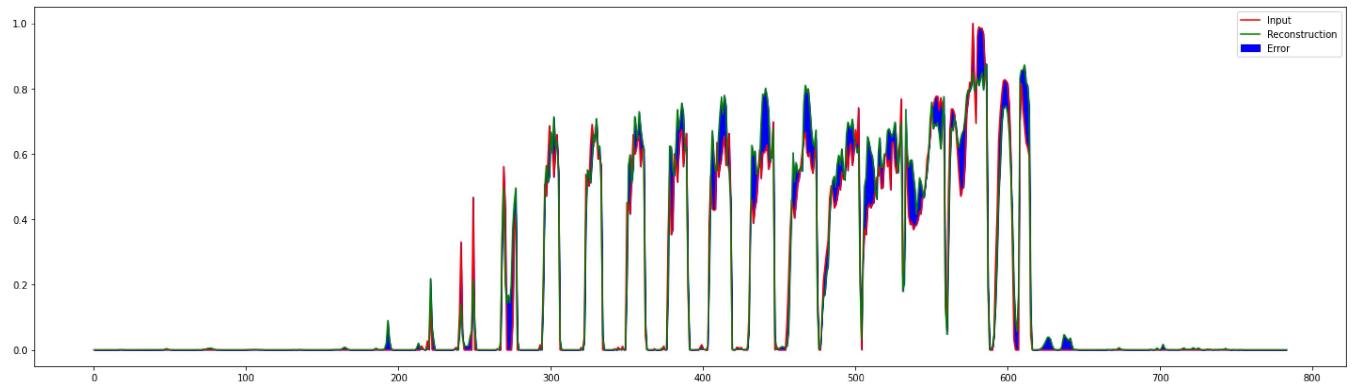
```
show_reconstructions(autoencoder, abnormal_test_data, 30)  
plt.show()
```



```
encoded_data = autoencoder.encoder(normal_test_data).numpy()  
decoded_data = autoencoder.decoder(encoded_data).numpy()
```

1-Dim plot of pixels of the first normal test data

```
plt.figure(figsize=(25,7))
plt.plot(normal_test_data[0].ravel(), 'r')
plt.plot(decoded_data[0].ravel(), 'g')
plt.fill_between(np.arange(28*28), decoded_data[0].ravel(), normal_test_data[0].ravel(), colo
plt.legend(labels=["Input", "Reconstruction", "Error"])
plt.show()
```



```
encoded_abn_data = autoencoder.encoder(abnormal_test_data).numpy()
decoded_abn_data = autoencoder.decoder(encoded_abn_data).numpy()
```

1-Dim plot of pixels of the first abnormal test data

```
plt.figure(figsize=(25,7))
plt.plot(abnormal_test_data[0].ravel(), 'r')
plt.plot(decoded_abn_data[0].ravel(), 'g')
plt.fill_between(np.arange(28*28), decoded_abn_data[0].ravel(), abnormal_test_data[0].ravel())
plt.legend(labels=["Input", "Reconstruction", "Error"])
plt.show()
```

▼ Distributions of the reconstruction losses and the calculation of the threshold.

0.6

Input

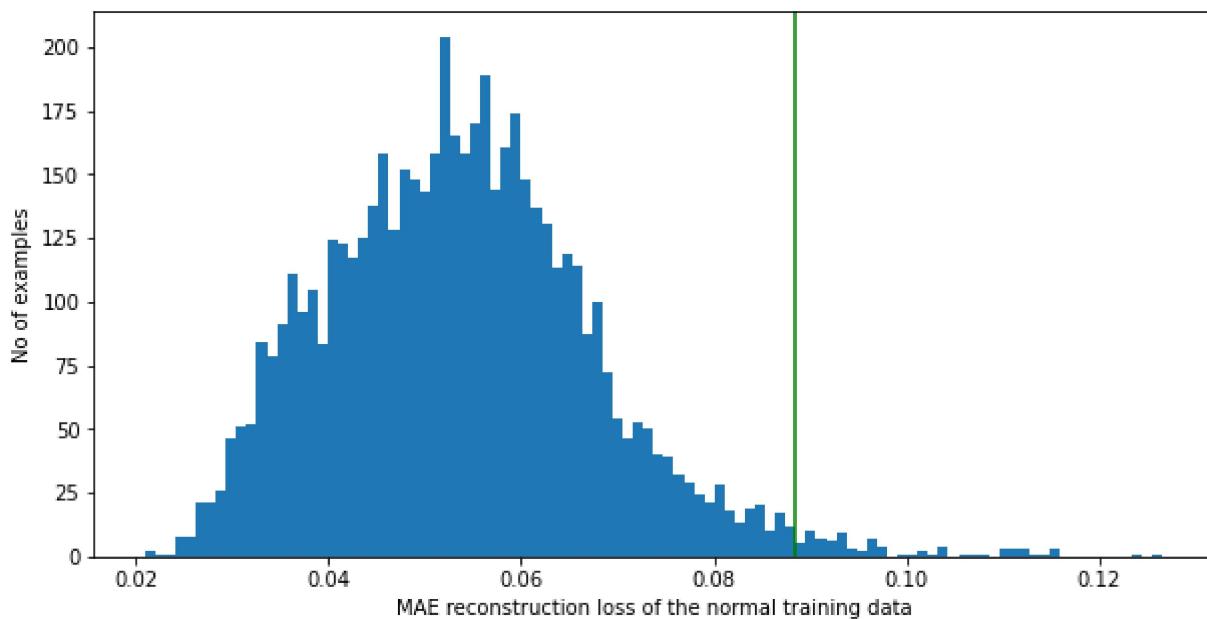
Distribution of the reconstruction losses of the normal training data

| Input

```

reconstructions = autoencoder.predict(normal_data)
train_loss = tf.keras.losses.mae(reconstructions.reshape(-1, 784), normal_data.reshape(-1, 784))
plt.figure(figsize=(10,5))
plt.hist(train_loss[None,:], bins=100)
threshold1 = np.mean(train_loss) + 2.5*np.std(train_loss)
plt.axvline(threshold1,c='g')
plt.xlabel("MAE reconstruction loss of the normal training data")
plt.ylabel("No of examples")
plt.show()

```



```

print("Mean: ", np.mean(train_loss))
print("Std: ", np.std(train_loss))

```

```

Mean:  0.05376415
Std:  0.013848975

```

```

threshold_train_mean_2_5_std = np.mean(train_loss) + 2.5*np.std(train_loss)
print("Threshold based on the mean of the training data MAE reconstruction losses + 2.5 std:")

```

```
Threshold based on the mean of the training data MAE reconstruction losses + 2.5 std: 0.085
```

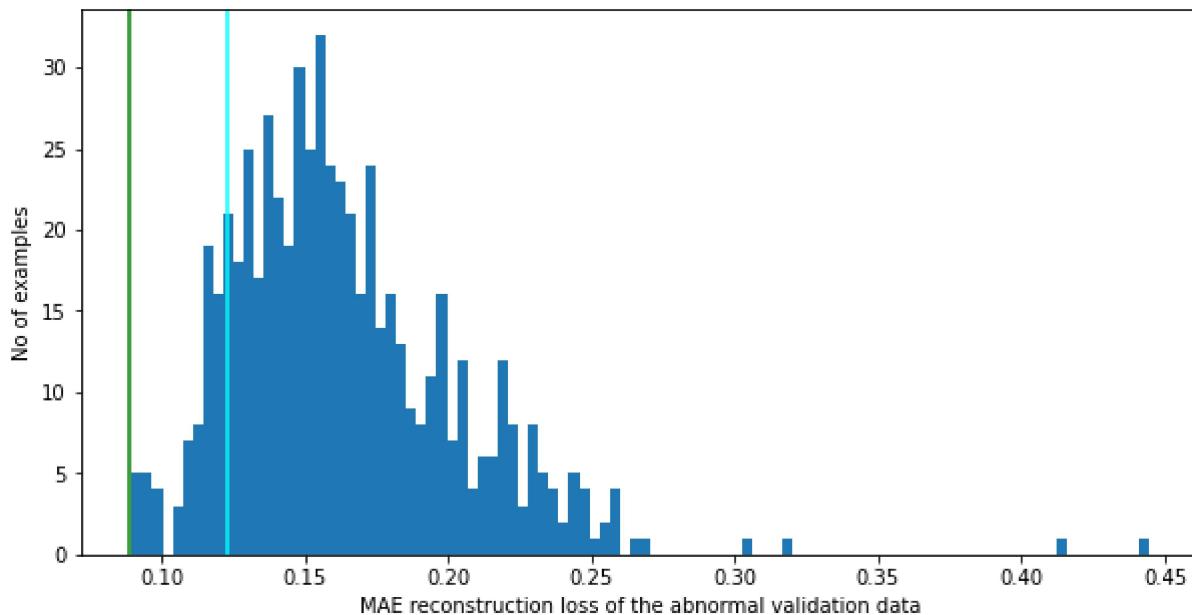
```
threshold1 = threshold_train_mean_2_5_std
```

Distribution of the reconstruction losses of the abnormal validation data

```

reconstructions = autoencoder.predict(abnormal_valid_data)
abn_valid_loss = tf.keras.losses.mae(reconstructions.reshape(-1,784), abnormal_valid_data.reshape(-1,784))
plt.figure(figsize=(10,5))
plt.hist(abn_valid_loss[None, :], bins=100)
threshold2 = np.mean(abn_valid_loss) - np.std(abn_valid_loss)
plt.axvline(threshold2,c='cyan')
plt.axvline(threshold1,c='g')
plt.xlabel("MAE reconstruction loss of the abnormal validation data")
plt.ylabel("No of examples")
plt.show()

```



```
abnormal_valid_mean_loss = np.mean(abn_valid_loss)
```

```
abnormal_valid_mean_loss , np.std(abn_valid_loss)
```

```
(0.16337438, 0.0400931)
```

```
threshold2 = abnormal_valid_mean_loss - np.std(abn_valid_loss)
print("Threshold2: ", threshold2)
```

```
Threshold2: 0.12328128
```

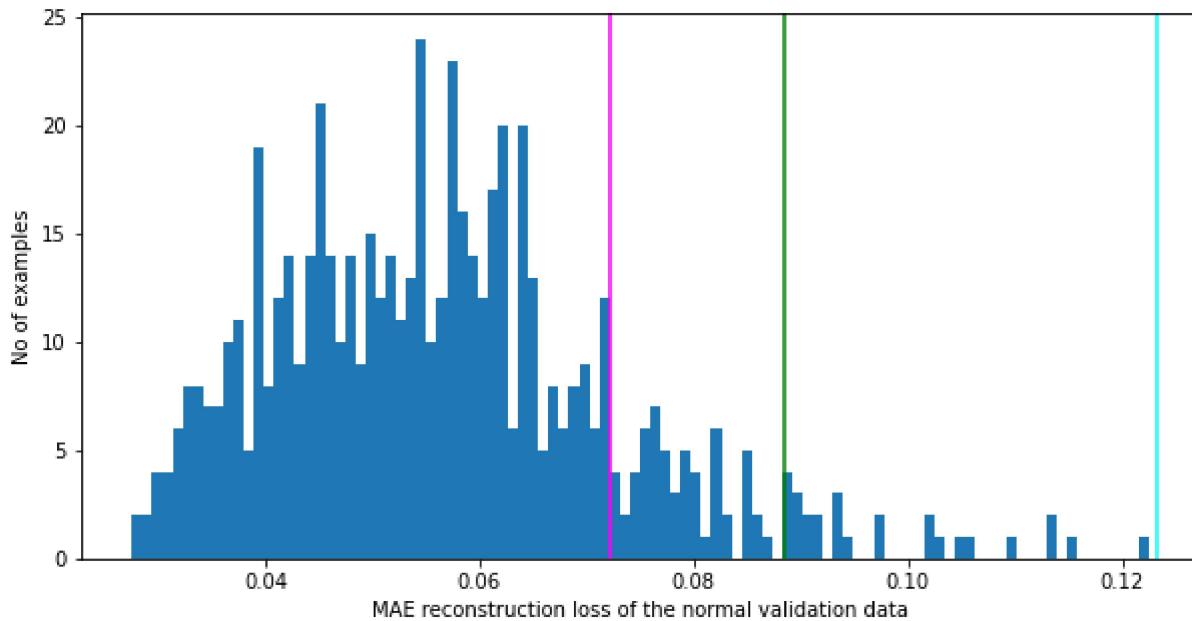
Distribution of the reconstruction losses of the normal validation data

```

reconstructions = autoencoder.predict(normal_valid_data)
nl_valid_loss = tf.keras.losses.mae(reconstructions.reshape(-1,784), normal_valid_data.reshape(-1,784))

```

```
plt.figure(figsize=(10,5))
plt.hist(nl_valid_loss[None, :], bins=100)
threshold3 = np.mean(nl_valid_loss) + np.std(nl_valid_loss)
plt.axvline(threshold3, c='magenta')
plt.axvline(threshold2, c='cyan')
plt.axvline(threshold1, c='g')
plt.xlabel("MAE reconstruction loss of the normal validation data")
plt.ylabel("No of examples")
plt.show()
```



```
normal_valid_mean_loss = np.mean(nl_valid_loss)
```

```
normal_valid_mean_loss , np.std(nl_valid_loss)
```

```
(0.05631686, 0.015880305)
```

```
threshold3 = normal_valid_mean_loss + np.std(nl_valid_loss)
print("Threshold3: ", threshold3)
```

```
Threshold3: 0.07219717
```

Calculation of a preliminary threshold based on $(\text{threshold2} + \text{threshold3}) / 2$ = Average of (mean + std of the distribution of the reconstruction losses of the normal validation data) and (mean - std of the distribution of the reconstruction losses of the abnormal validation data)

```
Avg_of_threshold_2_3 = (threshold2 + threshold3)/2
print("Average of threshold 2 and 3: ", Avg_of_threshold_2_3)
```

```
Average of threshold 2 and 3: 0.09773921966552734
```

```
threshold4 = Avg_of_threshold_2_3
```

Calculation of the threshold that gives the best accuracy on the validation data and set this as the threshold.

```
def predict(model, data, threshold):
    reconstructions = model.predict(data)
    loss = tf.keras.losses.mae(reconstructions.reshape(-1, 784), data.reshape(-1, 784))
    return tf.math.less(loss, threshold)
```

Double-click (or enter) to edit

```
increment = (abnormal_valid_mean_loss - normal_valid_mean_loss)/100
thresholds = np.arange(normal_valid_mean_loss, abnormal_valid_mean_loss, increment)
thrs_size = thresholds.shape[0]
accuracies = np.zeros(thrs_size)
for i in range(thrs_size):
    preds = predict(autoencoder, valid_data, thresholds[i])
    accuracies[i] = accuracy_score(preds, valid_labels_T_F)
argmax = np.argmax(accuracies)
valid_data_best_threshold = thresholds[argmax]
print("The best threshold based on validation data: ", valid_data_best_threshold)
```

The best threshold based on validation data: 0.09485756695270536

```
thr_acc = np.zeros((thrs_size, 2))
thr_acc[:, 0] = thresholds
thr_acc[:, 1] = accuracies
thr_acc[argmax-2:argmax+3]

array([[0.09271642, 0.98257261],
       [0.09378699, 0.98340249],
       [0.09485757, 0.98506224],
       [0.09592814, 0.98174274],
       [0.09699872, 0.98174274]])
```

```
threshold5 = valid_data_best_threshold
```

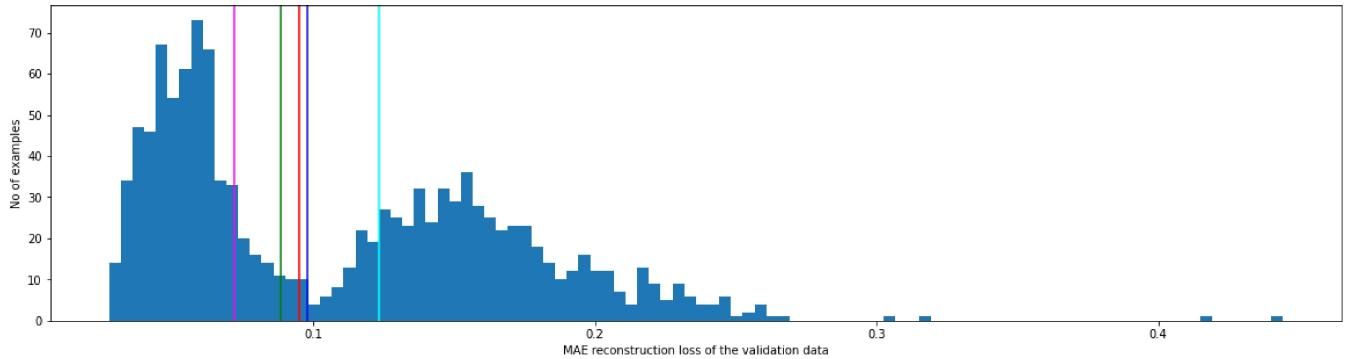
```
threshold = threshold5
```

Distribution of the reconstruction losses of all the validation data (normal and abnormal)

The blue line is threshold4 (= the average of threshold3 [magenta] and threshold2 [cyan]).

The red line is the threshold that gives the best accuracy for the validation data.

```
reconstructions = autoencoder.predict(valid_data)
valid_loss = tf.keras.losses.mae(reconstructions.reshape(-1,784), valid_data.reshape(-1,784))
plt.figure(figsize=(20,5))
plt.hist(valid_loss[None, :], bins=100)
plt.axvline(threshold, c='r')
plt.axvline(threshold4, c='b')
plt.axvline(threshold2, c='cyan')
plt.axvline(threshold3, c='magenta')
plt.xlabel("MAE reconstruction loss of the validation data")
plt.ylabel("No of examples")
plt.show()
```



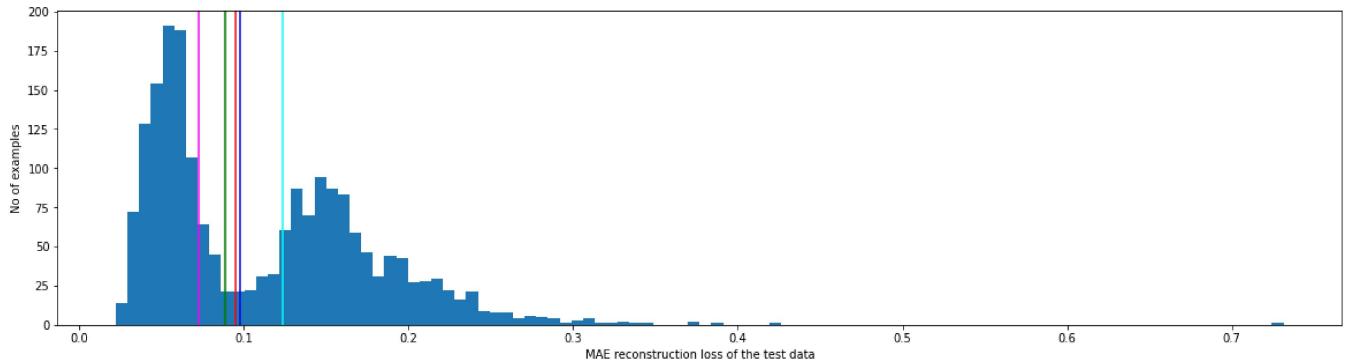
▼ Distribution of the reconstruction losses of the test data (normal and abnormal)

The blue line is threshold4 (= the average of threshold3 [magenta] and threshold2 [cyan]).

The red line is the threshold that gives the best accuracy for the validation data.

```
reconstructions = autoencoder.predict(test_data)
test_loss = tf.keras.losses.mae(reconstructions.reshape(-1,784), test_data.reshape(-1,784))
plt.figure(figsize=(20,5))
plt.hist(test_loss[None, :], bins=100)
plt.axvline(threshold, c='r')
plt.axvline(threshold4, c='b')
plt.axvline(threshold2, c='cyan')
plt.axvline(threshold3, c='magenta')
plt.xlabel("MAE reconstruction loss of the test data")
```

```
plt.ylabel("No of examples")
plt.show()
```



```
reconstructions = autoencoder.predict(normal_test_data)
nl_test_loss = tf.keras.losses.mae(reconstructions.reshape(-1,784), normal_test_data.reshape(
np.mean(nl_test_loss) , np.std(nl_test_loss))

(0.056795873, 0.018115388)
```

```
reconstructions = autoencoder.predict(abnormal_test_data)
abn_test_loss = tf.keras.losses.mae(reconstructions.reshape(-1,784), abnormal_test_data.reshape(
np.mean(abn_test_loss) , np.std(abn_test_loss)

(0.1684254, 0.048379235)
```

Calculation of the accuracy and the confusion matrix on the test data with threshold set based on the best threshold from the validation data

```
# def predict(model, data, threshold):
#     reconstructions = model.predict(data)
#     loss = tf.keras.losses.mae(reconstructions.reshape(-1, 784), data.reshape(-1, 784))
#     return tf.math.less(loss, threshold)

def print_stats(predictions, labels):
    cf = confusion_matrix(labels, predictions)
    print("Confusion Matrix: \n prediction: F      T ")
    print("          {}  {}".format(preds[preds == False].shape[0], preds[preds == True].sh
    print(" label: F  [[{}  {}]  {}]".format(cf[0,0], cf[0,1], test_labels_T_F[test_labels_T_
    print("          T  [[{}  {}]]  {}]".format(cf[1,0], cf[1,1], test_labels_T_F[test_labels_T_
    print("Accuracy = {}".format(accuracy_score(labels, predictions)))
    print("Normal Test Data Mean = {}".format(np.mean(nl_test_loss)))
    print("Normal Test Data Standard Deviation = {}".format(np.std(nl_test_loss)))
    print("Abnormal Test Data Mean = {}".format(np.mean(abn_test_loss)))
    print("Abnormal Test Data Standard Deviation = {}".format(np.std(abn_test_loss)))
    print("Precision = {}".format(precision_score(labels, predictions)))
    print("Recall = {}".format(recall_score(labels, predictions)))
```

```

print(accuracy_score(labels, predictions))
print(np.mean(nl_test_loss))
print(np.std(nl_test_loss))
print(np.mean(abn_test_loss))
print(np.std(abn_test_loss))
print(precision_score(labels, predictions))
print(recall_score(labels, predictions))
print(accuracy_score(labels, predictions), np.mean(nl_test_loss), np.std(nl_test_loss), np.
precision_score(labels, predictions), recall_score(labels, predictions))

```

```

preds = predict(autoencoder, test_data, valid_data_best_threshold)
print_stats(preds, test_labels_T_F)

```

↳ Confusion Matrix:

		F	T
		1013	987
label: F	F	[990 10]	1000
	T	[23 977]	1000

Accuracy = 0.9835

Normal Test Data Mean = 0.0567958727478981

Normal Test Data Standard Deviation = 0.018115388229489326

Abnormal Test Data Mean = 0.16842539608478546

Abnormal Test Data Standard Deviation = 0.048379234969615936

Precision = 0.9898682877406282

Recall = 0.977

0.9835

0.056795873

0.018115388

0.1684254

0.048379235

0.9898682877406282

0.977

0.9835 0.056795873 0.018115388 0.1684254 0.048379235 0.9898682877406282 0.977

```

print("Threshold =", valid_data_best_threshold)

```

Threshold = 0.09485756695270536

```

print(confusion_matrix(test_labels_T_F, preds))

```

```

[[990 10]
 [ 23 977]]

```

Extra accuracy info

Just informative. Please record the above accuracy.

Accuracy on the test data with threshold set based on (threshold2 + threshold3) / 2 = Average of (mean + std of the distribution of the reconstruction losses of the normal validation data) and (mean - std of the distribution of the reconstruction losses of the abnormal validation data)

```
preds = predict(autoencoder, test_data, Avg_of_threshold_2_3)
print_stats(preds, test_labels_T_F)
```

Confusion Matrix:

		F	T
		1006	994
label: F	[[987	13]	1000
	T	[19	981]]

Accuracy = 0.984

Normal Test Data Mean = 0.0567958727478981

Normal Test Data Standard Deviation = 0.018115388229489326

Abnormal Test Data Mean = 0.16842539608478546

Abnormal Test Data Standard Deviation = 0.048379234969615936

Precision = 0.9869215291750503

Recall = 0.981

0.984

0.056795873

0.018115388

0.1684254

0.048379235

0.9869215291750503

0.981

0.984 0.056795873 0.018115388 0.1684254 0.048379235 0.9869215291750503 0.981

Accuracy on the test data with threshold set based on the mean of the training data MAE reconstruction losses + 2.5 std

```
preds = predict(autoencoder, test_data, threshold_train_mean_2_5_std)
print_stats(preds, test_labels_T_F)
```

Confusion Matrix:

		F	T
		1030	970
label: F	[[996	4]	1000
	T	[34	966]]

Accuracy = 0.981

Normal Test Data Mean = 0.0567958727478981

Normal Test Data Standard Deviation = 0.018115388229489326

Abnormal Test Data Mean = 0.16842539608478546

Abnormal Test Data Standard Deviation = 0.048379234969615936

Precision = 0.9958762886597938

Recall = 0.966

0.981

0.056795873

0.018115388

0.1684254

0.048379235

0.9958762886597938

0.966

0.981 0.056795873 0.018115388 0.1684254 0.048379235 0.9958762886597938 0.966

✓ 0s completed at 3:25 PM

